

BISO560R Spring 2014 Homework 2

Pittard

Due by 11:59 PM on February 14, 2014

Instructions

Send responses in a plain text file name LastName_Firstname_HW2.txt or LastName_Firstname_HW2.R. You can use RStudio to create the latter. We will run your commands at the R console to verify the statements. Email to BOTH dvandom@emory.edu and wsp@emory.edu

1 Data Frames - 30 points

Write a function called `my.sampler` that fulfills the following requirements:
It will take a data frame, such as `mtcars`, and “split” it up based on a given grouping variable. Then it will process each group and sample a specified number of records, (without replacement), from each group. When it is finished processing all groups it will combine those results into a data frame and return it. You need only accommodate a single group/factor per function call. As an example:

```
my.sampler(mtcars,mtcars$cyl,3)
```

	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
Toyota Corona	21.5	4	120.1	97	3.70	2.465	20.01	1	0	3	1
Honda Civic	30.4	4	75.7	52	4.93	1.615	18.52	1	1	4	2
Merc 230	22.8	4	140.8	95	3.92	3.150	22.90	1	0	4	2
Merc 280	19.2	6	167.6	123	3.92	3.440	18.30	1	0	4	4
Mazda RX4	21.0	6	160.0	110	3.90	2.620	16.46	0	1	4	4
Merc 280C	17.8	6	167.6	123	3.92	3.440	18.90	1	0	4	4
Hornet Sportabout	18.7	8	360.0	175	3.15	3.440	17.02	0	0	3	2
AMC Javelin	15.2	8	304.0	150	3.15	3.435	17.30	0	0	3	2
Cadillac Fleetwood	10.4	8	472.0	205	2.93	5.250	17.98	0	0	3	4


```
my.sampler(ChickWeight,ChickWeight$Diet,2) # ChickWeight is built-in to R
```

	weight	Time	Chick	Diet
20	138	14	2	1
57	197	16	5	1
253	145	16	23	2

270	49	2	25	2
361	221	16	32	3
389	41	0	35	3
513	135	12	45	4
554	322	21	48	4

Your function will have to work with the other potential grouping variables from the given data set (in the case of `mtcars` - `gear`, `transmission`, `carb`, and `vs`). Once you have split the data frame on a factor you should be able to, as part of your looping or apply logic, determine how many rows are in each category using functions you've learned about in class. We will also run your function on the `ChickWeight` data and the `diamonds` dataset, (in the `ggplot2` package), to insure that there is nothing specific in it to `mtcars`

THINGS TO CHECK: Insure that the number of unique values that the grouping variable takes on is less than 10. So if you are given a factor to split on then check that it assumes at most 10 unique values - otherwise complain to the user with an appropriate error message.

You should also do error checking to make sure that the number of records you are attempting to sample is not larger than the available number of records in the group. In that case then your program will stop with an informative message. For example in looking at the `mtcars` data, one can see that there is only one car that has a value of 6 in the `carb` column, thus

```
my.sampler(mtcars,mtcars$carb,2)  would fail:
```

```
Error in my.sampler(mtcars, mtcars$carb, 2) :
  Number of requested samples 2 exceeds available records
```

As would the following because MPG really isn't a grouping variable:

```
my.sampler(mtcars,mtcars$mpg,5)
Error in my.sampler(mtcars, mtcars$mpg, 5) :
  Grouping factor has 25 unique values. Must be less than 10
```

Table 1: Arguments to `my.sampler` function

Argument	Purpose
<code>my.df</code>	a dataframe (e.g. <code>mtcars</code> , <code>ChickWeight</code> , etc)
<code>my.group</code>	a grouping variable from <code>my.df</code>
<code>numtosample</code>	the number of records to sample from each value of <code>my.group</code>

2 DNA - 30 points

2.1 seqgen - 10 points

Write a function called `seqgen` that returns a randomly generated DNA nucleotide sequence. Here we will restrict ourselves to using four characters - ACGT. For this exercise all 4 bases occur with the same probability. Your function should take the following arguments:

Table 2: Arguments to `seqgen` function

Argument	Purpose
<code>size</code>	the number of characters in the desired sequence
<code>collapsed</code> (default TRUE)	return a collapsed character string or, if FALSE, a vector
<code>upper.case</code> (default TRUE)	return upper or lower case

As an example, a collapsed string would look like "AGTTGG" whereas an uncollapsed string would look like "A" "G" "T" "T" "G" "G". Your function should validate the size of the requested sequence to make sure it makes sense. That is check for negative numbers or character strings. The second and third arguments should have defaults of TRUE and TRUE respectively. Here are some examples of how the function might be called:

```
seqgen(20,T,F)
[1] "ggtataacaatccccgtgggt"
```

```
seqgen(20,T,T)
[1] "GCTGCAATCTCGGCCTTGAT"
```

```
seqgen(20,F,T)
[1] "G" "T" "A" "A" "G" "C" "A" "T" "C" "A" "G" "A" "A" "C" "G" "C" "T" "C" "C" "C"
```

```
seqgen(0)
Error in seqgen(0) : Sorry - size needs to be numeric and greater than or equal to 1
```

2.2 revcomp - 20 points

Write a companion function called `revcomp` that, given any valid sequence generated by `seqgen`, will return the complimentary sequence. *Note that for this assignment use of the "chartr" function is forbidden.* In DNA terminology certain bases have specific compliments. The character A changes to T, C changes to G, G changes to C, and T changes to A. As an example the string ACGTTGA would have a compliment of: TGCAACT.

Moreover, if we wanted to take the “reverse” compliment of a string we would first reverse the given string and then apply the compliment substitution as indicated. So given the same string as above, “ACGTTGA”, we would first reverse it to get “AGTTGCA” after which we would apply the compliment substitutions to get a reverse compliment of “TCAACGT”. Use http://www.bioinformatics.org/sms/rev_comp.html to check your work. Make sure to use the drop down menu to reflect whatever it is you are testing: “complement”, “reverse compliment”, etc. Your function will take the following arguments:

Table 3: Arguments to revcomp function

Argument	Purpose
sequence	A sequence string like that returned by seqgen
reverse (default FALSE)	If TRUE then return the compliment of the reverse DNA string
collapsed (default FALSE)	If TRUE then return a collapsed string

Here are some examples of how revcomp might be called:

```
my.seq = seqgen(20)
```

```
my.seq
[1] "C" "A" "T" "T" "A" "T" "C" "A" "T" "C" "G" "C" "A" "A" "T" "T" "C" "T" "T"
[20] "A"
```

```
revcomp(my.seq)
[1] "g" "t" "a" "a" "t" "a" "g" "t" "a" "g" "c" "g" "t" "t" "a" "a" "g" "a" "a"
[20] "t"
```

```
revcomp(seqgen(20),collapsed=T)
[1] "ccactcctattgatattata"
```

If you want a specific test sequence then use the following. However there should be nothing in your code that is specific to this string:

```
TAACGGCGGCAGGCGTATGG
```

```
revcomp("TAACGGCGGCAGGCGTATGG", reverse= F, collapsed=T )
[1] "attgccgccgtccgcatacc"
```

```
revcomp("TAACGGCGGCAGGCGTATGG", reverse=T, collapsed=T)
[1] "ccatacgcctgccgccgtta"
```

3 Quartile - 40 points

Write a function called “quartile.nist” which computes the requested percentiles of a given vector `x`. The calling sequence should look like this:

```
quartile.nist(x, probs=c(0.25,0.50,0.75))
```

Table 4: Arguments to quartile.nist function

Argument	Purpose
<code>x</code>	a numeric vector of any size
<code>probs</code>	a vector of desired percentile(s) where ($0 < \text{probs} < 1$)

The default value for `probs` should be `c(.25,.50,.75)`. Check for values in `probs` that are less than or equal to zero or greater than or equal to 1. For hints on how to structure this function see the psuedocode given below in Figure 1. Here are some examples of how this function can be used:

```
x = c(19,15.5,15.0,20.5,18.3,20.9,11.7,24.3,23.9)
quartile.nist(x)
```

```
 25%  50%  75%
15.3 19.0 22.4
```

```
quartile.nist(x,c(.40,.60))
 40%  60%
18.3 20.5
```

```
set.seed(145)
testx = rnorm(1000)
quartile.nist(testx)
 25%    50%    75%
-0.5810  0.0756  0.7543
```

3.1 Background

Suppose that we have a data vector x_1, x_2, \dots, x_n . The p th percentile is conceptually meant to be a value Q_p such that p proportion of the observations fall below Q_p . There are different ways to compute the percentiles / quartiles. Here we will use the “n+1” method. Consider the data set:

```
Xn = 19.0 15.5 15.0 20.5 18.3 20.9 11.7 24.3 23.9
```

```
x = c(19,15.5,15.0,20.5,18.3,20.9,11.7,24.3,23.9)
```

Here $n = 9$. We will first need to order these values:

```
sortedx = 11.7 15.0 15.5 18.3 19.0 20.5 20.9 23.9 24.3
```

To compute the 20th percentile of the data we calculate Q_{20} as follows:

$$(n + 1) * p = (9 + 1) * .20 = 2$$

So the value representing the 20th percentile is `sortedx[2]` which is 15. Thus for computed Q values that wind up being an integer then you simply use the Q value to index into the sorted vector and you are done. Now, let's compute the 33rd percentile. Here we must use interpolation which will require some adjustment to our approach.

$$Q = (n + 1) * p = (9 + 1) * 0.33 = 3.3$$

Note that we have a number with a fractional part. HINT: It will be very useful for you to code up some logic to break up this number into its whole component (3) and its fractional component (0.3). You will need these values early on in your function to do comparisons. So here we take the third observation (the whole part of Q) of the SORTED vector and add to it 0.3 (the fractional part of Q) times the distance between it and the next highest (4th) observation. So:

$$Q_{33} = \text{sortedx}[3] + 0.3 * (\text{sortedx}[4] - \text{sortedx}[3]) \text{ is } 16.3$$

```
quartile.nist(x,.33)
33%
16.34
```

So that is the approach you implement for computed Q values that have a fractional part. Well almost. There are "boundary conditions" that we must consider. For low or high values of prob (like less than .10 or greater than .90) you need to pay attention to the computed Q value. Consider the case where prob is 0.95. $Q = (9 + 1) * .95 = 9.5$. Since there are only 9 elements in the original vector you can't interpolate it since there are no other elements above element 9. So in this case you would subtract 1 from the whole part (giving 8) and then apply the given interpolation formula:

$$Q_{95} = \text{sortedx}[8] + 0.5 * (\text{sortedx}[9] - \text{sortedx}[8]) = 24.1$$

```
quartile.nist(x,.95)
95%
24.1
```

Also consider the case where prob is 0.05. $Q = (9 + 1) * .05 = 0.5$. Note that there is no *zero-th* element of the vector. So we add 1 to Q and apply the given interpolation formula. So Q is 0.5 but we add 1 to it to get 1.5 and then apply the formula:

```
Q05 = sortedx[1] + 0.5*(sortedx[2] - sortedx[1]) = 13.3
```

```
quartile.nist(x,.05)
```

```
5%
```

```
13.35
```

Figure 1: Possible Pseudocode for nist.quartile function

```
determine the length, (n), of the input vector x
sort the vector x for later use
setup an empty vector that will be used to return the requested quartiles

for each value in the probabilities vector do
  compute Q using the (n+1)*p formula
  split up Q into its whole and fractional parts
  if Q has a non zero fractional part
    implement some logic to check for low and hi boundary conditions;
    apply interpolation formula;
  else
    Q is a whole number with a zero fractional part;
    Use whole part of Q to index into the sorted array;
  end
  append the computed Q to the retrun vector
end
name the return vector elements to reflect the requested percentiles (e.g. 25%, 50%, etc)
```