

Aggregation

Data Aggregation: The process of taking some data and putting it into a form that lends itself easily to summary, e.g. replace groups of observations with summary statistics.

Data Restructuring: Change the structure of the data so that its new form is more convenient for a specific purpose (usually analysis).

R has many ways to do either. Before you start writing code to do some of these activities check around to see what functions exist. Chances are there are some good tools available. We will explore many of them in this session.

Aggregation

So all of this is basically a lesson that there are functions available in R that can help with aggregation.

Command	Purpose
table, xtabs	Summarize grouping variables with grouping variables - Create Contingency Tables
tapply	Summarize a continuous variable by grouping variables
aggregate	Summarize continuous variable(s) by grouping variables
melt, cast	Reshape data and summarize continuous variable(s) by grouping variables

Aggregation - tables

One of the most basic forms of aggregation is to put things into tables for easy counting. This is usually done with categorical variables to obtain "count" information. We can then do things like Chi-Square tests.

```
letters[1:10]                # Built in letters char vector
[1] "a" "b" "c" "d" "e" "f" "g" "h" "i" "j"

# sample 20 times from the 1st 10 letters with replacement
set.seed(222)
(my.sample = sample(letters[1:10],20,replace=TRUE) )
[1] "j" "a" "e" "a" "j" "j" "d" "e" "f" "b" "e" "a" "b" "h" "j" "a" "g" "a"
"b"
[20] "c"

table(my.sample)
my.sample
a b c d e f g h j
5 3 1 1 3 1 1 1 4
```

Aggregation - tables

```
prop.table(table(my.sample))
```

```
my.sample
```

a	b	c	d	e	f	g	h	j
0.25	0.15	0.05	0.05	0.15	0.05	0.05	0.05	0.20

```
sum(prop.table(table(my.sample)))
```

```
[1] 1
```

Aggregation - tables

```
set.seed(123)
coinsamp = sample(c('H','T'),1000,T)

table(coinsamp)
coinsamp
  H    T
507 493
chisq.test(table(coinsamp),p=c(0.5,0.5))
```

Chi-squared test for given probabilities

```
data:  table(coinsamp)
X-squared = 0.196, df = 1, p-value = 0.658
```

Aggregation - tables

```
set.seed(123)
coinsamp = sample(c('H','T'),1000,T,prob=c(.70,.30))
chisq.test(table(coinsamp),p=c(0.5,0.5))
```

Chi-squared test for given probabilities

```
data:  table(coinsamp)
X-squared = 168.1, df = 1, p-value < 2.2e-16
```

Aggregation - tables

One of the most basic forms of aggregation is to put things into tables for easy counting. This is usually done with categorical variables to obtain "count" information. We can then do things like Chi-Square tests.

```
table(mtcars$am)
```

```
 0  1  
19 13
```

```
table(mtcars$cyl)
```

```
 4  6  8  
11  7 14
```

Aggregation - tables

However, we usually create tables out of a data frame that contains some categories or factors. Using the popular mtcars data set let's find out how many 4,6,8 and cylinder cars there are for each category of transmission (auto or manual).

```
table(mtcars$am,mtcars$cyl)
```

	4	6	8
0	3	4	12
1	8	3	2

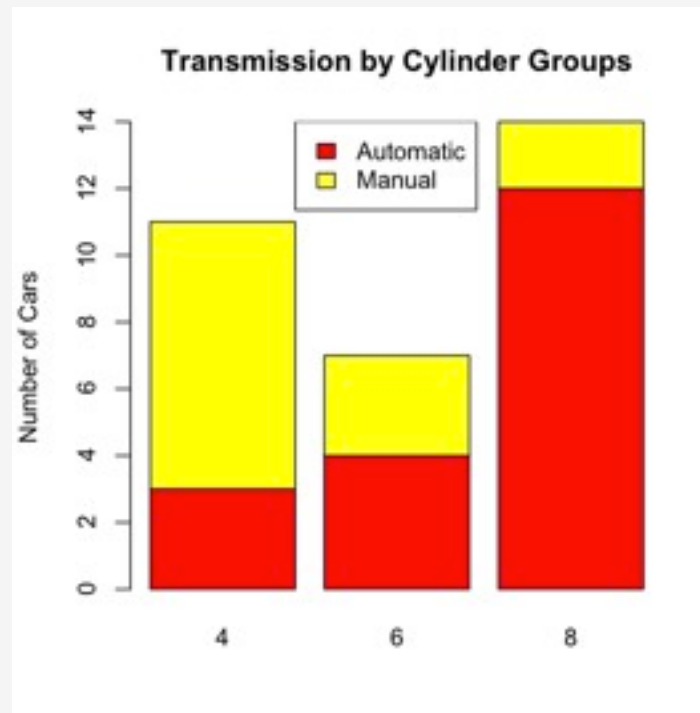
```
table(Trans=mtcars$am,Cyl=mtcars$cyl)
```

	Cyl		
Trans	4	6	8
0	3	4	12
1	8	3	2

Aggregation - tables

```
mytable = table(mtcars$am, mtcars$cyl)
barplot(mytable, main="Transmission by Cylinder Groups",
        ylab="Number of Cars",col=heat.colors(2))

legend(1.2,14,fill=heat.colors(2),c("Automatic","Manual"))
```



Aggregation - tables

```
mtcars = transform(mtcars, am = factor(am,labels=c("Automatic","Manual")))
```

```
my.table = table(mtcars$am,mtcars$cyl)
```

	4	6	8
Automatic	3	4	12
Manual	8	3	2

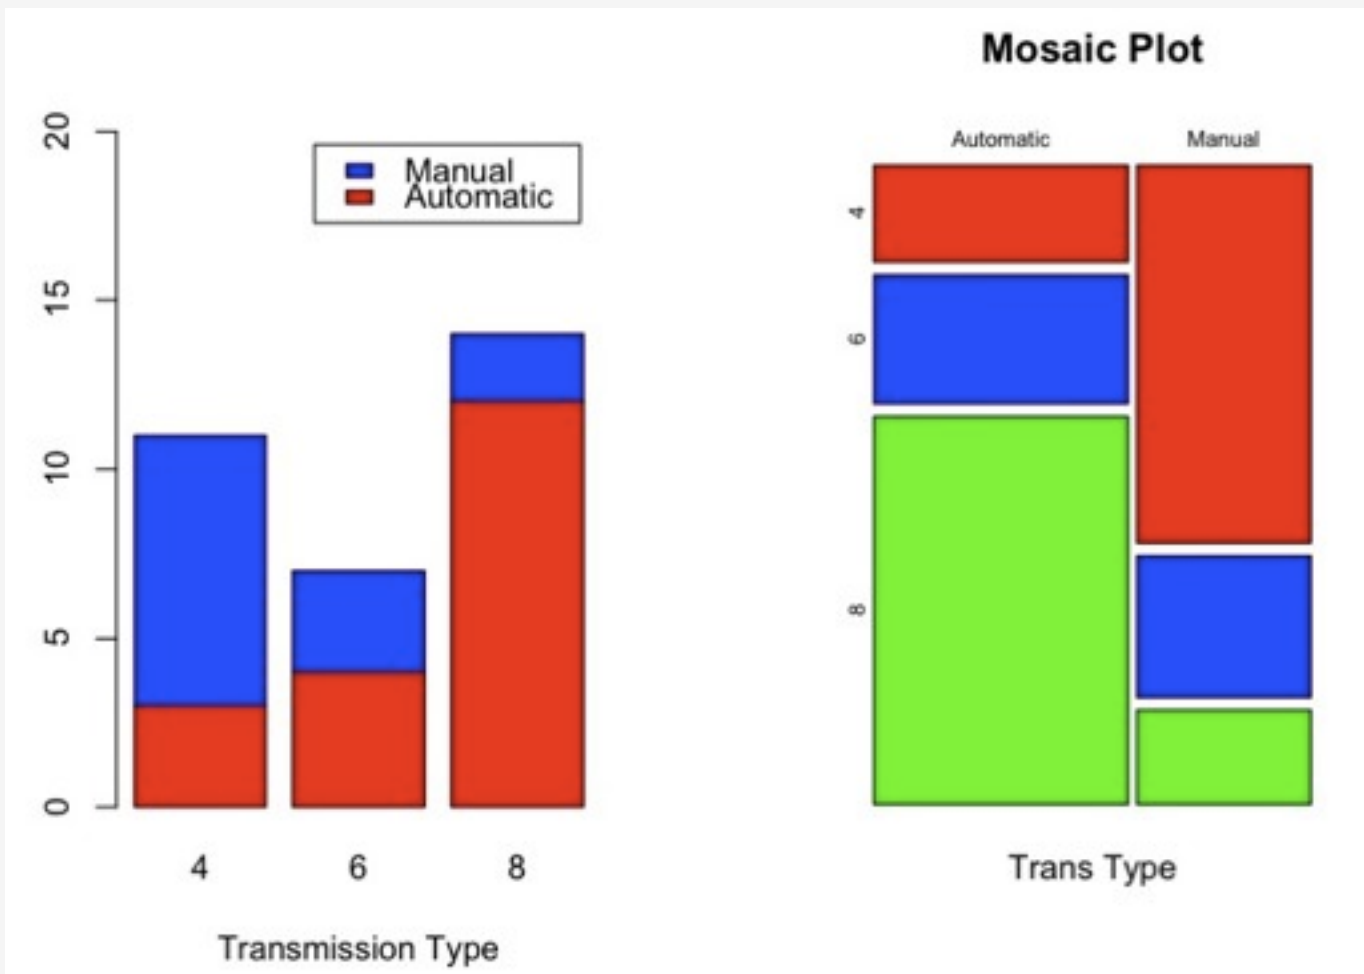
```
par(mfrow=c(1,2))
```

```
barplot(my.table,legend=T,col=c("red","blue","green"),  
        xlab="Transmission Type",ylim=c(0,20))
```

```
# Here we get a Mosaic Plot
```

```
plot(my.table,col=c("red","blue","green"),main="Mosaic Plot",  
     xlab="Trans Type", ylab="Number of Gears")
```

Aggregation - tables



Aggregation - tables

The table function can work on really big data frames. Remember the large Chicago crime data frame ?

```
url = "http://www.bimcore.emory.edu/BIOS560R/DATA.DIR/chi\_crimes.csv"
chi = read.table(url,header=T,sep=",")
nrow(chi)
[1] 334141

crime.table = table(chi$Arrest,chi$Ward)
```

Aggregation - tables

```
names(chi)
```

[1] "Case.Number"	"ID"
[3] "Date"	"Block"
[5] "IUCR"	"Primary.Type"
[7] "Description"	"Location.Description"
[9] "Arrest"	"Domestic"
[11] "Beat"	"District"
[13] "Ward"	"FBI.Code"
[15] "X.Coordinate"	"Community.Area"
[17] "Y.Coordinate"	"Year"
[19] "Latitude"	"Updated.On"
[21] "Longitude"	"Location"

Aggregation - tables

```
crime.table = table(chi$Arrest,chi$District)
```

	1	2	3	4	5	6
false	8811	10319	12806	15422	10974	13464
true	3296	3129	4843	4367	4284	5768

	7	8	9	10	11	12
false	13877	16422	12140	10359	12821	6858
true	6273	5964	4516	4657	8977	1916

	13	14	15	16	17	18
false	5624	9980	8704	8868	7728	11008
true	1460	2557	5681	1885	1945	3170

	19	20	22	24	25	31
false	12378	4333	8342	7070	14099	0
true	3230	1341	2403	2428	5559	6

Aggregation - tables

The table function can work on really big data frames. Remember the large Chicago crime data frame ?

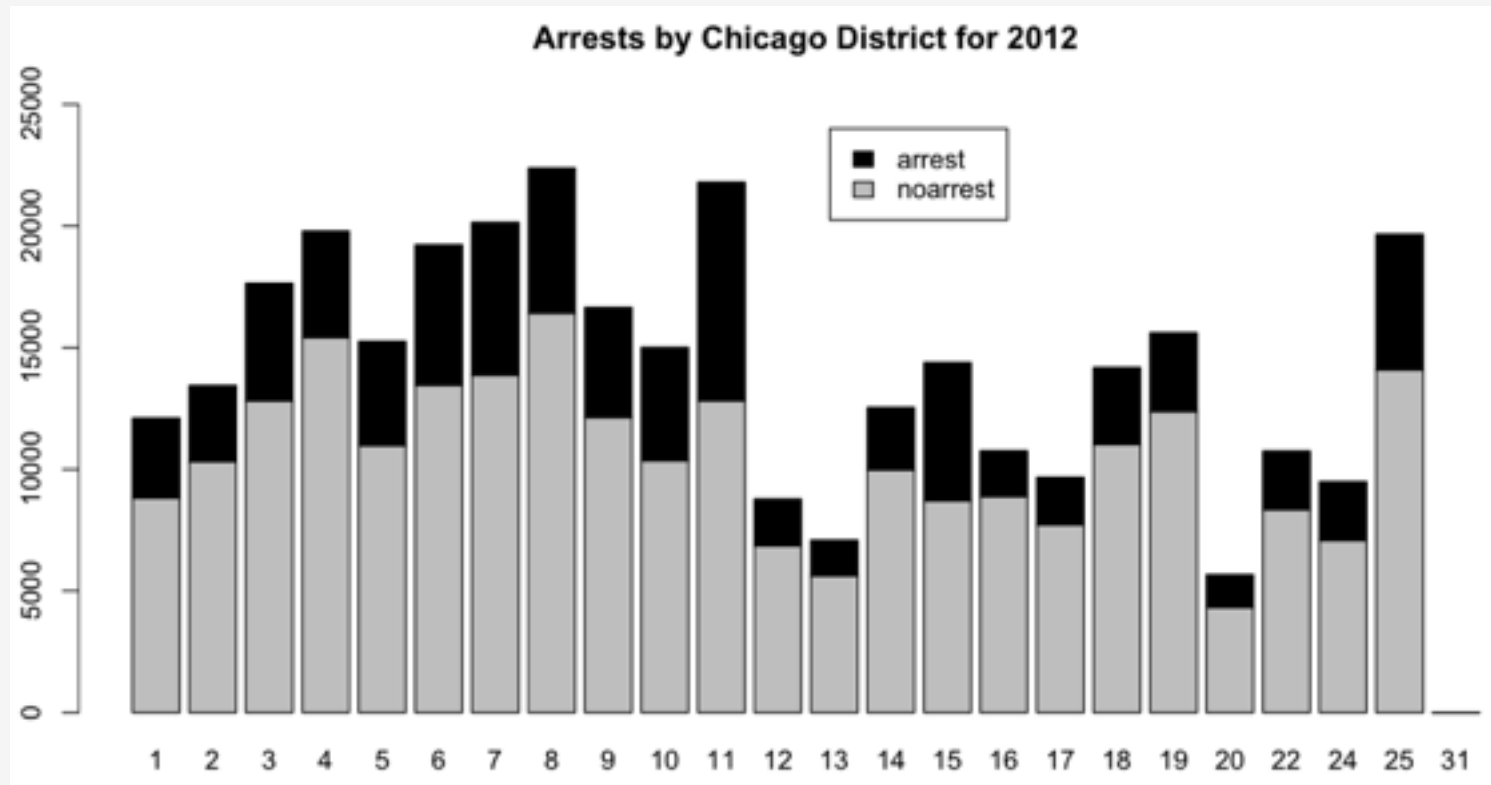
```
url = "http://www.bimcore.emory.edu/BIOS560R/DATA.DIR/chi\_crimes.csv"
chi = read.table(url,header=T,sep=",")
nrow(chi)
[1] 334141
crime.table = table(chi$Arrest,chi$District)

barplot(crime.table,col=c("gray","black"),
        ylim=c(0,25000),
        main="Arrests by Chicago District for 2012")

legend(15,24000,fill=c("black","grey"),c("arrest","noarrest"))
```

Aggregation - tables

The table function can work on really big data frames. Remember the large Chicago crime data frame ?



Aggregation - tables

Remember that we can take continuous variables and turn them into categories.

```
my.cut = cut(mtcars$mpg,breaks=quantile(mtcars$mpg),
             labels=c("Bad","Not Good","Good","Great"),include.lowest=T)
```

```
my.cut
```

```
[1] Good      Good      Good      Good      Not Good
[6] Not Good  Bad       Great     Good      Not Good
[11] Not Good  Not Good  Not Good  Bad       Bad
[16] Bad       Bad       Great     Great     Great
[21] Good      Not Good  Bad       Bad       Not Good
[26] Great     Great     Great     Not Good  Good
[31] Bad       Good
Levels: Bad Not Good Good Great
```

```
table(MPG=my.cut,Cyl=mtcars$cyl)
```


	Cyl		
MPG	4	6	8
Bad	0	0	8
Not Good	0	3	6
Good	4	4	0
Great	7	0	0

Aggregation - tables

```
plot(table(MPG=my.cut,mtcars$cyl),col=terrain.colors(3),  
      main="Gas Rating by Cylinder Group")
```

Aggregation - tables

Cyl				
MPG	4	6	8	
Bad	0	0	8	
Not Good	0	3	6	
Good	4	4	0	
Great	7	0	0	



Aggregation - margins

Its also a very common thing to want margins for the rows or columns or maybe even both. Once we have a table we have some flexibility.

```
mtcars = transform(mtcars, am = factor(am,labels=c("Automatic","Manual")))
```

```
my.table = table(mtcars$am,mtcars$cyl)
```

	4	6	8
Automatic	3	4	12
Manual	8	3	2

<http://www.cyclismo.org/tutorial/R/tables.html#data>

Aggregation - margins

But this is too hard. Let's look in R to see if there is a function that will already do this for us.

```
> ??margins
base::apply          Apply Functions Over Array Margins
base::margin.table   Compute table margin
base::prop.table     Express Table Entries as Fraction of Marginal Table
..
..

AND A BUNCH OF OTHER STUFF
..

stats::addmargins    Puts Arbitrary Margins on Multidimensional Tables or Arrays
```

Aggregation - margins

```
addmargins(my.table,1)      # Place the sums of each column as a new row
```

	4	6	8
Automatic	3	4	12
Manual	8	3	2
Sum	11	7	14

```
addmargins(my.table,2)      # Place sums of each row as a new column
```

	4	6	8	Sum
Automatic	3	4	12	19
Manual	8	3	2	13

```
addmargins(my.table,c(1,2)) # Get Both
```

	4	6	8	Sum
Automatic	3	4	12	19
Manual	8	3	2	13
Sum	11	7	14	32

Aggregation - margins

But remember that we found some interesting functions before:

Help files with alias or concept or title matching 'margins' using fuzzy matching:

<code>base::apply</code>	Apply Functions Over Array Margins
<code>base::margin.table</code>	Compute table margin
<code>base::prop.table</code>	Express Table Entries as Fraction of Marginal Table
<code>car::mmps</code>	Marginal Model Plotting
<code>gbm::plot.gbm</code>	Marginal plots of fitted gbm objects

Aggregation - margins

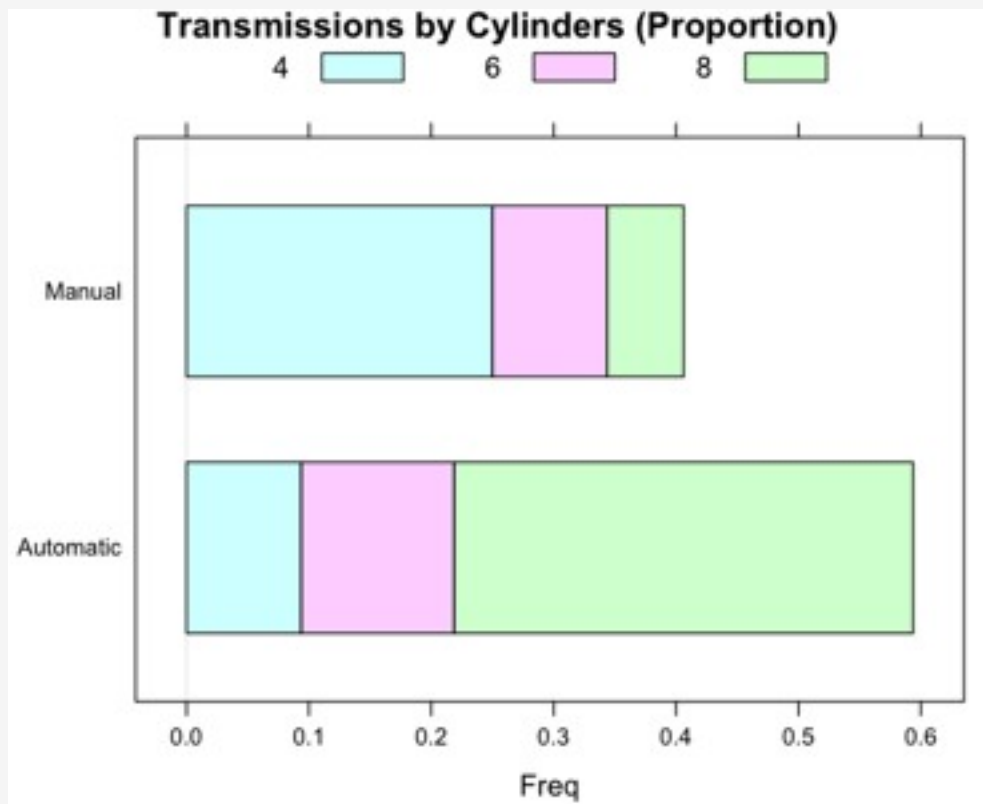
But there is an easier way. Use the `prop.table` command, which, if you look at the help pages for it, is simply an alias to the `sweep` command without all the confusing arguments.

```
prop.table(my.table,1)
      4      6      8
Automatic 0.1578947 0.2105263 0.6315789
Manual    0.6153846 0.2307692 0.1538462

rowSums(prop.table(my.table,1)) # The rows add up to one as they should
Automatic      Manual
      1      1
```


Aggregation - margins

```
library(lattice)
barchart(prop.table(my.table), auto.key=list(columns=3),
          main="Transmissions by Cylinders (Proportion)")
```



Aggregation - 3D Tables

You should also note that tables can be 3-dimensional though this can get tricky when viewing.

```
my.3d.table = table(mtcars$am,mtcars$gear,mtcars$cyl)
```

```
, , = 4
```

	3	4	5
Automatic	1	2	0
Manual	0	6	2

```
, , = 6
```

	3	4	5
Automatic	2	2	0
Manual	0	2	1

```
, , = 8
```

	3	4	5
Automatic	12	0	0
Manual	0	0	2

Aggregation - 3D Tables

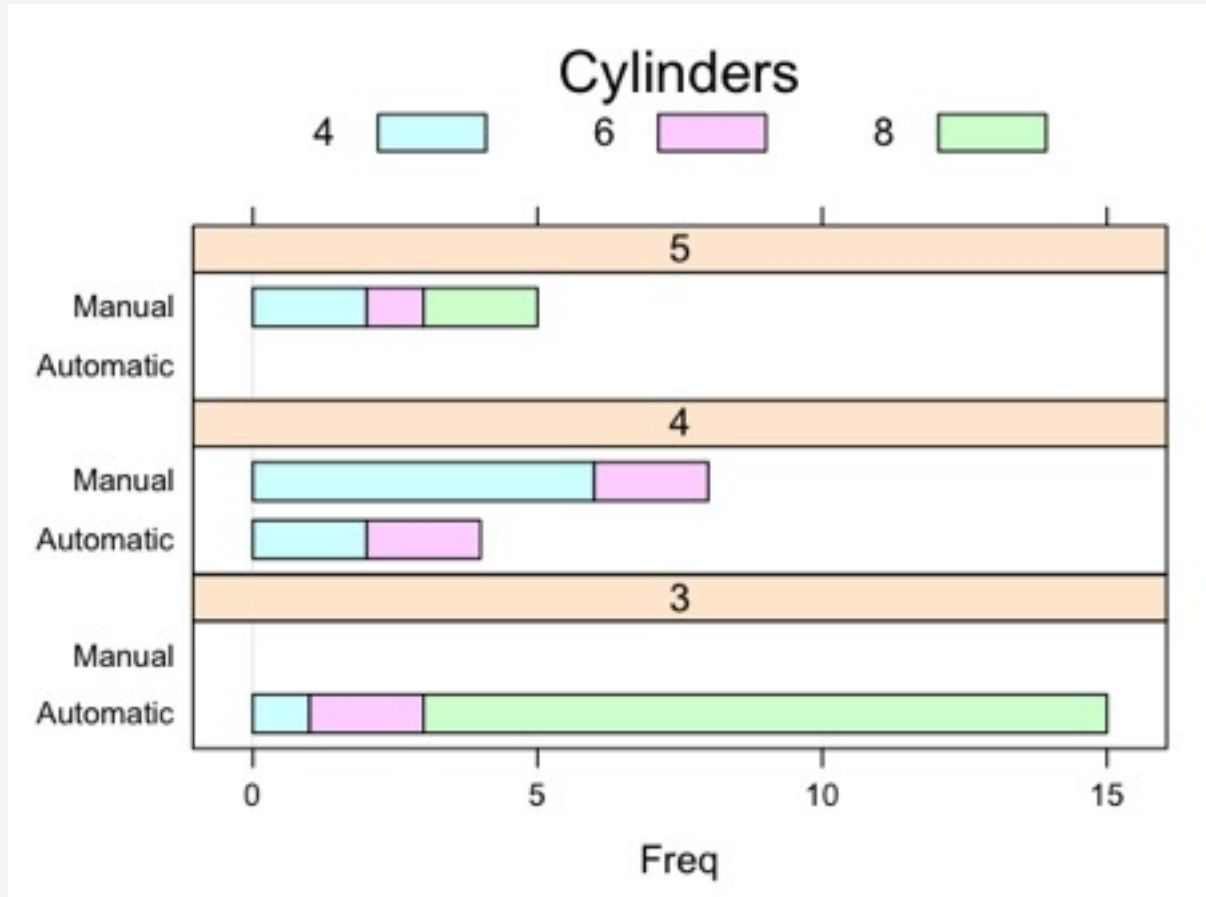
There is a function called `fable` that flattens your table for ease of viewing:

```
fable(my.3d.table)
      4  6  8
Automatic 3   1  2 12
          4   2  2  0
          5   0  0  0
Manual    3   0  0  0
          4   6  2  0
          5   2  1  2
>
```

Aggregation - 3D Tables

The lattice graphics package can help here:

```
barchart(my.3d.table,layout=c(1,3),auto.key=list(title="Cylinders",columns=3))
```



Aggregation - xtabs

This wouldn't be R unless there were some other function that does pretty much the same thing as table. There is another function called "xtabs". Some people feel that it is superior to the table function because it supports a formula interface when creating tables. The formula interface requires the "+" symbol to be on the right of any grouping variable.

```
xtabs(~am,mtcars)      # Transmission type by number of cylinders
```

```
am
```

```
  0  1
```

```
19 13
```

```
# Number in each cylinder group
```

```
xtabs(~cyl,mtcars)
```

```
cyl
```

```
  4  6  8
```

```
11  7 14
```

Aggregation - xtabs

This wouldn't be R unless there were some other function that does pretty much the same thing as table. There is another function called "xtabs". Some people feel that it is superior to the table function because it supports a formula interface when creating tables. The formula interface requires the "+" symbol to be on the right of any grouping variable.

```
xtabs(~am + cyl, mtcars)      # Transmission type by number of cylinders
```

```
      cyl
am
  Automatic 3  4 12
  Manual    8  3  2
```

```
xtabs(~cyl + am, mtcars)     # Number of cylinders by transmission type
```

```
      am
cyl Automatic Manual
  4          3      8
  6          4      3
  8         12      2
```

Aggregation - xtabs

A three dimensional table is possible

```
xtabs(~gear + cyl + am, mtcars)  
, , am = Automatic
```

	cyl		
gear	4	6	8
3	1	2	12
4	2	2	0
5	0	0	0

```
, , am = Manual
```

	cyl		
gear	4	6	8
3	0	0	0
4	6	2	0
5	2	1	2

Aggregation - xtabs

Part of the reason we might use xtabs is because the formula interface is also used when testing for mutual independence. If we create a table with xtabs then we use the same formula in the loglm function.

```
mytab = xtabs(~gear + cyl + am, mtcars)
```

```
library(MASS)
loglm(~gear + cyl + am, mytab)
```

Call:

```
loglm(formula = ~gear + cyl + am, data = mytab)
```

Statistics:

	X^2	df	P(> X^2)
Likelihood Ratio	51.94755	12	6.333779e-07
Pearson	49.24950	12	1.891657e-06

Aggregation - xtabs

Also, xtabs allows us to use a subset argument to pull out certain records:

```
xtabs(~gear + cyl, mtcars)
```

```
      cyl  
gear  4  6  8  
  3  1  2 12  
  4  8  4  0  
  5  2  1  2
```

```
xtabs(~gear + cyl, mtcars, subset=mpg < 20)
```

```
      cyl  
gear  6  8  
  3  1 12  
  4  2  0  
  5  1  2
```

Aggregation - tapply

Summarizing a continuous variable in terms of one or more grouping variables

Aggregation - tapply

Up until now we've been working with tables created from categorical variables. Even in our example that used a continuous variable we wound up using the `cut` command to group the continuous variables into categories.

We can easily summarize continuous variables in terms of a grouping variable using the **tapply** and **by** commands.

continuous factor summary function

```
tapply(mtcars$mpg, mtcars$am, mean)
```

```
Automatic    Manual  
17.14737    24.39231
```

```
tapply(mtcars$mpg, list(Transmission=mtcars$am, Cyl=mtcars$cyl), mean)
```

```
          Cyl  
Transmission    4    6    8  
          0 22.9 19.1 15.1  
          1 28.1 20.6 15.4
```

Aggregation - tapply

One limitation of the tapply command is that we can't summarize more than one continuous variable at a time.

This works:

```
tapply(mtcars$mpg, list(mtcars$am, mtcars$vs), mean)
      0      1
Automatic 0.8008991 0.9339769
Manual    1.6366124 1.7982418
```

But this does not:

```
tapply(list(mtcars$mpg,mtcars$hp), list(mtcars$am, mtcars$cyl), mean)
Error in tapply(list(mtcars$mpg, mtcars$hp), list(mtcars$am, mtcars$cyl), :
  arguments must have same length
```

Aggregation - The aggregate command

This is the more general command for aggregation of continuous data in terms of some factors or categories:

```
aggregate(mtcars['mpg'],list(Transmission=mtcars$am),mean)
```

	Transmission	mpg
1	Automatic	17.14737
2	Manual	24.39231

```
aggregate(mtcars[c('mpg', 'hp')],list(Transmission_Type=mtcars$am),mean)
```

	Transmission_Type	mpg	hp
1	Automatic	17.14737	160.2632
2	Manual	24.39231	126.8462

Aggregation - The aggregate command

```
aggregate(mtcars[c('mpg','hp')],  
          list(Transmission_Type = mtcars$am,  
               Cylinders = mtcars$cyl), mean)
```

	Transmission_Type	Cylinders	mpg	hp
1	Automatic	4	22.90000	84.66667
2	Manual	4	28.07500	81.87500
3	Automatic	6	19.12500	115.25000
4	Manual	6	20.56667	131.66667
5	Automatic	8	15.05000	194.16667
6	Manual	8	15.40000	299.50000

Aggregation - The aggregate command

The aggregate command also has a formula interface if you find that to be more convenient. Many do since it gives you an argument to specify the data frame you are trying to summarize. This saves typing.

```
aggregate(mpg ~ am, mtcars, mean)
```

	am	mpg
1	auto	17.14737
2	manual	24.39231

```
aggregate(mpg ~ am + cyl, mtcars, mean)
```

	am	cyl	mpg
1	auto	4	22.90000
2	manual	4	28.07500
3	auto	6	19.12500
4	manual	6	20.56667
5	auto	8	15.05000
6	manual	8	15.40000

Aggregation - The aggregate command

The aggregate command also has a formula interface if you find that to be more convenient. Many do since it also gives you an argument to specify the data frame you are trying to summarize. This saves typing.

```
aggregate( cbind(mpg,wt,hp) ~ am + cyl, mtcars, mean)
```

	am	cyl	mpg	wt	hp
1	Automatic	4	22.9	2.94	84.7
2	Manual	4	28.1	2.04	81.9
3	Automatic	6	19.1	3.39	115.2
4	Manual	6	20.6	2.75	131.7
5	Automatic	8	15.1	4.10	194.2
6	Manual	8	15.4	3.37	299.5

Aggregation - The aggregate command

So here is how this would look using the older bracket style:

```
aggregate( mtcars[c('mpg','wt','hp')],  
          list(trans=mtcars$am, cyl=mtcars$cyl), mean)
```

	trans	cyl	mpg	wt	hp
1	Automatic	4	22.9	2.94	84.7
2	Manual	4	28.1	2.04	81.9
3	Automatic	6	19.1	3.39	115.2
4	Manual	6	20.6	2.75	131.7
5	Automatic	8	15.1	4.10	194.2
6	Manual	8	15.4	3.37	299.5

Aggregation - The aggregate command

```
my.snps = read.table("http://www.bimcore.emory.edu/nsnps.csv",header=F,sep=" ")
names(my.snps)=c("X1","X2","SNP1","SNP2")
new.snps = cbind(my.snps,co=sample(c("case","control"),30,TRUE))
```

```
head(new.snps)
      X1   X2 SNP1 SNP2      co
1 -0.737 2.77   AA   BB control
2 -1.839 5.70   AA   BB   case
3 -0.470 4.89   AA   BB   case
4  0.851 6.38   AA   BB control
5 -1.690 8.11   AA   BB   case
6 -1.480 4.33   AA   BB control
```

```
# Let's summarize the mean of X1 and X2 across the "co" column
```

```
aggregate(cbind(X1,X2)~co,new.snps,mean)
      co      X1      X2
1  case -0.0371  4.77
2 control -0.0163  5.15
```

Aggregation - The aggregate command

```
head(new.snps)
      X1   X2 SNP1 SNP2   co
1 -0.737 2.77   AA   BB control
2 -1.839 5.70   AA   BB   case
3 -0.470 4.89   AA   BB   case
4  0.851 6.38   AA   BB control
5 -1.690 8.11   AA   BB   case
6 -1.480 4.33   AA   BB control
```

```
# Let's summarize the mean of X1 and X2 across the SNP1 column
```

```
aggregate(cbind(X1,X2)~SNP1,new.snps,mean)
      SNP1      X1      X2
1     aa  0.222  4.88
2     Aa  0.279  4.98
3     AA -0.573  5.19
```

Aggregation - The melt and cast commands

- * melt and cast belong to a package called "reshape" which is an add on package for doing data frame manipulation and , if desired, aggregation.
- * Some feel that these commands are superior to the built in commands such as aggregate, tapply, and by. They do provide a general way to break down data into a general format and then reassemble it so we discuss it here.
- * Let's look at a simple data set. You will need to tell melt which variables are "id" variables, (usually factors and/or characters), and which are "measured" variables, (continuous variables that are measurements on the id variables).

Aggregation - The melt and cast commands

First we need to install the reshape package:

```
install.packages("reshape")  
library(reshape)
```

Second we need to understand the difference between "long" format and "wide" format.

```
data(smiths)  
smiths
```

	subject	time	age	weight	height
1	John Smith	1	33	90	1.87
2	Mary Smith	1	NA	NA	1.54

This is "wide" format - all the data relating to an observation is on one line only

Aggregation - The melt and cast commands

Wide format is the most convenient for humans. Everything is on one line only.

```
smiths
  subject time age weight height
1 John Smith    1  33     90   1.87
2 Mary Smith    1  NA     NA   1.54
```

Long format is better for doing certain statistical tests such as ANOVA

```
  subject variable value
1 John Smith      time  1.00
2 Mary Smith      time  1.00
3 John Smith      age 33.00
4 Mary Smith      age  NA
5 John Smith     weight 90.00
6 Mary Smith     weight  NA
7 John Smith     height  1.87
8 Mary Smith     height  1.54
```

Aggregation - The melt and cast commands

We need ways to conveniently move back and forth between the two. The reshape package includes two commands to help with this: melt and cast

```
melt(smiths)
Using subject as id variables
  subject variable value
1 John Smith      time  1.00
2 Mary Smith      time  1.00
3 John Smith      age   33.00
4 Mary Smith      age    NA
5 John Smith     weight 90.00
6 Mary Smith     weight  NA
7 John Smith     height  1.87
8 Mary Smith     height  1.54
```

Melt can figure out what the appropriate id variable is - in this case "Subject". We could specify this explicitly:

```
melt(smiths, id.var="subject")
```

Aggregation - The melt and cast commands

The other important argument to melt is "measure.var". this tells melt what variables you want to use for measuring and or analysis.

```
melt(smiths,measure.var=c("time","age","weight","height"))
```

	subject	variable	value
1	John Smith	time	1.00
2	Mary Smith	time	1.00
3	John Smith	age	33.00
4	Mary Smith	age	NA
5	John Smith	weight	90.00
6	Mary Smith	weight	NA
7	John Smith	height	1.87
8	Mary Smith	height	1.54

Aggregation - The melt and cast commands

```
my.df = data.frame(id=paste("P",1:4,sep=""),
                    age=c(30.0,36.0,43.0,29.0),
                    pulse=c(78.2,74.4,82.3,91.9),
                    height=c(72,68,69,71),
                    group=c("ctrl","ctrl","treat","treat"))
```

```
my.df
  id age pulse height group
1 P1  30  78.2     72  ctrl
2 P2  36  74.4     68  ctrl
3 P3  43  82.3     69 treat
4 P4  29  91.9     71 treat
```

```
sapply(my.df,class)
      id      age      pulse      height      group
"factor" "numeric" "numeric" "numeric"  "factor"
```

If we don't give melt any information then it will use "id" and "group" as identification variables.

Aggregation - The melt and cast commands

```
> sapply(my.df, class)
      id      age      pulse      height      group
"factor" "numeric" "numeric" "numeric"  "factor"
```

```
> melt(my.df)
```

Using id, group as id variables

	id	group	variable	value
1	P1	ctrl	age	30.0
2	P2	ctrl	age	36.0
3	P3	treat	age	43.0
4	P4	treat	age	29.0
5	P1	ctrl	pulse	78.2
6	P2	ctrl	pulse	74.4
7	P3	treat	pulse	82.3
8	P4	treat	pulse	91.9
9	P1	ctrl	height	72.0
10	P2	ctrl	height	68.0
11	P3	treat	height	69.0
12	P4	treat	height	71.0

```
my.melt = melt(my.df)
```

Aggregation - The melt and cast commands

```
melt(my.df)

# Would be the same as

melt(my.df, id.vars = c("id","group"))

# Would be the same as

melt(my.df, id.vars = c("id","group"),
      measure.vars = c("age","pulse","height"))

# Would be the same as

melt(my.df, measure.vars = c("age","pulse","height"))
```

Aggregation - The melt and cast commands

```
cast(my.melt, group~variable, mean)
```

```
  group age pulse height  
1  ctrl  33  76.3     70  
2 treat  36  87.1     70
```

Aggregation - The melt and cast commands

```
cast(my.melt, group~variable, mean)
```

	group	age	pulse	height
1	ctrl	33	76.3	70
2	treat	36	87.1	70

```
cast(my.melt, group~variable, mean, subset=variable==c("height","age"))
```

	group	age	height
1	ctrl	36	72
2	treat	29	69

Aggregation - The melt and cast commands

```
url = "http://www.bimcore.emory.edu/BIOS560R/DATA.DIR/sixsmiths.csv"  
sixsmiths = read.table(url,header=T,sep=",")
```

```
sixsmiths  
  subject time age weight height  
1 John Smith   1  33     90   1.87  
2 Mary Smith   1  28     85   1.54  
3 John Smith   2  33     90   1.87  
4 Mary Smith   2  28     89   1.54  
5 John Smith   3  33     90   1.87  
6 Mary Smith   3  28     92   1.54
```

Aggregation - The melt and cast commands

```
my.melt = melt(sixsmiths,id.vars=c("subject","time"))
```

	subject	time	variable	value
1	John Smith	1	age	33.00
2	Mary Smith	1	age	28.00
3	John Smith	2	age	33.00
4	Mary Smith	2	age	28.00
5	John Smith	3	age	33.00
6	Mary Smith	3	age	28.00
7	John Smith	1	weight	90.00
8	Mary Smith	1	weight	85.00
9	John Smith	2	weight	90.00
10	Mary Smith	2	weight	89.00
11	John Smith	3	weight	90.00
12	Mary Smith	3	weight	92.00
13	John Smith	1	height	1.87
14	Mary Smith	1	height	1.54
15	John Smith	2	height	1.87
16	Mary Smith	2	height	1.54
17	John Smith	3	height	1.87
18	Mary Smith	3	height	1.54

Aggregation - The melt and cast commands

```
cast(my.melt,time~variable,mean)
```

	time	age	weight	height
1	1	30.5	87.5	1.71
2	2	30.5	89.5	1.71
3	3	30.5	91.0	1.71

```
cast(my.melt,time~variable,mean,subset=variable==c("weight"))
```

	time	weight
1	1	87.5
2	2	89.5
3	3	91.0

Aggregation - Summary

So all of this is basically a lesson that there are functions available in R that can help with aggregation.

Command	Purpose
table, xtabs	Create Contingency Tables
tapply, by, split	Summarize a continuous variable by a grouping variables
aggregate	Summarize continuous variable(s) by grouping variables
melt, cast	Reshape data and summarize continuous variable(s) by grouping variables

Aggregation - Summary

So what are the rules here ?

If you want to look at count based data across one or more categories you can use tables or the xtabs functions.

If you want to look at count based data across one or more categories and you anticipate doing lots of stat tests on them then use the xtabs function.

If you want to look at one continuous variable summarized in some way across multiple categories then use tapply or by.

If you want to look at continuous variable(s) summarized in some way across multiple categories then use aggregate or melt/cast.