

Matrices - Doing Calculations

Let's look at some examples involving calculations on matrices:

```
set.seed(123)
```

```
X = matrix(rpois(9,1.5),nrow=3)
```

```
colnames(X) = c("aspirin","paracetamol","nurofen")
```

```
rownames(X) = paste("Trial",1:3,sep=".")
```

```
X
```

| | aspirin | paracetamol | nurofen |
|---------|---------|-------------|---------|
| Trial.1 | 1 | 3 | 1 |
| Trial.2 | 2 | 4 | 3 |
| Trial.3 | 1 | 0 | 1 |

```
mean(X[,3]) # Mean of the 3rd column  
[1] 1.666667
```

```
var(X[3,]) # Variance of the 3rd row  
[1] 0.3333333
```

Matrices - Doing Calculations

Let's look at some examples involving calculations on matrices. But there are some general functions to help with this kind of thing:

```
X
      aspirin paracetamol nurofen
Trial.1      1          3       1
Trial.2      2          4       3
Trial.3      1          0       1
```

```
rowSums(X)
Trial.1 Trial.2 Trial.3
      5      9      2
```

```
colSums(X)
aspirin paracetamol nurofen
      4          7       5
```

Maybe columns represent protein expression and you are trying to determine if there are differences between the mean expression levels.

The R Book - Michael J. Crawley

Matrices - Doing Calculations

But there are some general functions to help with this kind of thing:

```
rowMeans(X)
  Trial.1 Trial.2 Trial.3
1.666667 3.000000 0.666667

colMeans(X)
  aspirin paracetamol nurofen
1.333333  2.333333  1.666667

colMeans(X)[3]
nurofen
1.666667
```

These are fast and can work on very large matrices. Though be careful if you have missing values in your data.

The R Book - Michael J. Crawley

Matrices - Doing Calculations - apply

Its worth pointing out that you can do similar things with the apply function. It allows you to plug in any function - not just the mean function.

```
X
      aspirin paracetamol nurofen
Trial.1      3          1        3
Trial.2      2          2        2
Trial.3      2          0        5

apply(X,1,summary)           # 1 is for rows
      Trial.1 Trial.2 Trial.3
Min.      1.000      2.0  0.0000
1st Qu.    1.000      2.5  0.5000
Median     1.000      3.0  1.0000
Mean       1.667      3.0  0.6667
3rd Qu.    2.000      3.5  1.0000
Max.       3.000      4.0  1.0000

apply(X,2,summary)           # 2 is for columns
      aspirin paracetamol nurofen
Min.      1.000      0.000  1.000
1st Qu.    1.000      1.500  1.000
Median     1.000      3.000  1.000
Mean       1.333      2.333  1.667
3rd Qu.    1.500      3.500  2.000
Max.       2.000      4.000  3.000
```

The R Book - Michael J. Crawley

Matrices - Doing Calculations - apply

Its worth pointing out that you can do similar things with the apply function. It allows you to plug in any function - not just the mean function.

```
apply(X,1,mean)
  Trial.1 Trial.2 Trial.3
2.000000 2.333333 0.333333
```

This is equivalent to:

```
rowMeans(X)
  Trial.1 Trial.2 Trial.3
2.000000 2.333333 0.333333
```

Let's "scale"/"center" the values in the rows. We subtract each value from the mean of its row

```
apply(X,1,function(x) (x-mean(x))) # Scale the values in the rows
```

```
      Trial.1 Trial.2 Trial.3
aspirin    -1  1.666667  0.666667
paracetamol  0 -0.333333 -0.333333
nurofen     1 -1.333333 -0.333333
```

Matrices - Doing Calculations - apply

Let's find what rows have values greater than 2. Let's also find the row that has the largest number of values greater than 2. Sound hard ? Not really.

```
X > 2
      aspirin paracetamol nurofen
Trial.1  FALSE         FALSE    TRUE
Trial.2   TRUE         FALSE   FALSE
Trial.3  FALSE         FALSE   FALSE

apply(X > 2, 1, sum)           # Its a tie it seems
Trial.1 Trial.2 Trial.3
      1      1      0

max(apply(X > 2, 1, sum))
[1] 1

which(apply(X > 2, 1, sum) == max(X>2) )
Trial.1 Trial.2
      1      2
```

This works because we can sum TRUE and FALSE values since R gives a value of "1" and "0" respectively.

```
as.numeric(TRUE)
[1] 1

as.numeric(FALSE)
[1] 0
```

Matrices - Linear Algebra

R supports common linear algebra operations also.

```
A = matrix(c(1,3,2,2,8,9),3,2)
```

```
A
      [,1] [,2]
[1,]    1    2
[2,]    3    8
[3,]    2    9
```

```
t(A)
      [,1] [,2] [,3]
[1,]    1    3    2
[2,]    2    8    9
```

$$\begin{bmatrix} 1 & 2 \\ 3 & 8 \\ 2 & 9 \end{bmatrix}^T = \begin{bmatrix} 1 & 3 & 2 \\ 2 & 8 & 9 \end{bmatrix}$$

<http://bendixcarstensen.com/APC/linalg-notes-BxC.pdf>

Matrices - Linear Algebra

A

```
      [,1] [,2]  
[1,]    1    2  
[2,]    3    8  
[3,]    2    9
```

B = matrix(c(5,8,4,2),2,2)

A %*% B

```
      [,1] [,2]  
[1,]   21    8  
[2,]   79   28  
[3,]   82   26
```

$$\begin{bmatrix} 1 & 2 \\ 3 & 8 \\ 2 & 9 \end{bmatrix} \begin{bmatrix} 5 & 4 \\ 8 & 2 \end{bmatrix} = \left[\begin{bmatrix} 1 & 2 \\ 3 & 8 \\ 2 & 9 \end{bmatrix} \begin{bmatrix} 5 \\ 8 \end{bmatrix} : \begin{bmatrix} 1 & 2 \\ 3 & 8 \\ 2 & 9 \end{bmatrix} \begin{bmatrix} 4 \\ 2 \end{bmatrix} \right]$$
$$= \begin{bmatrix} 1 \cdot 5 + 2 \cdot 8 & 1 \cdot 4 + 2 \cdot 2 \\ 3 \cdot 5 + 8 \cdot 8 & 3 \cdot 4 + 8 \cdot 2 \\ 2 \cdot 5 + 9 \cdot 8 & 2 \cdot 4 + 9 \cdot 2 \end{bmatrix} = \begin{bmatrix} 21 & 8 \\ 79 & 28 \\ 82 & 26 \end{bmatrix}$$

<http://bendixcarstensen.com/APC/linalg-notes-BxC.pdf>

Matrices - Linear Algebra

The inverse of a $n \times n$ matrix A is the matrix B (which is also $n \times n$) that when multiplied by A gives the identity matrix.

```
A = matrix(1:4,2,2)
```

A

```
      [,1] [,2]  
[1,]     1     3  
[2,]     2     4
```

```
B = solve(A)
```

B

```
      [,1] [,2]  
[1,]    -2  1.5  
[2,]     1 -0.5
```

```
A %*% B
```

We get the identity matrix

```
      [,1] [,2]  
[1,]     1     0  
[2,]     0     1
```

<http://bendixcarstensen.com/APC/linalg-notes-BxC.pdf>

Matrices - Linear Algebra

Suppose you have the following system of equations. This can be represented as:

$$\begin{array}{rcl} x_1 + 3x_2 & = & 7 \\ 2x_1 + 4x_2 & = & 10 \end{array}$$

$$\begin{bmatrix} 1 & 3 \\ 2 & 4 \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \end{bmatrix} = \begin{bmatrix} 7 \\ 10 \end{bmatrix} \text{ i.e. } Ax = b$$

```
A
    [,1] [,2]
[1,]    1    3
[2,]    2    4

b = c(7,10)
x = solve(A) %*% b
x
    [,1]
[1,]    1
[2,]    2
```

Since $A^{-1}A = I$ and since $Ix = x$ we have

$$x = A^{-1}b = \begin{bmatrix} -2 & 1.5 \\ 1 & -0.5 \end{bmatrix} \begin{bmatrix} 7 \\ 10 \end{bmatrix} = \begin{bmatrix} 1 \\ 2 \end{bmatrix}$$

<http://bendixcarstensen.com/APC/linalg-notes-BxC.pdf>

Lists - Intro

- * Lists address the situation where we need to store information of different types in a single structure.
- * Remember that vectors and matrices restrict us to only one data type at a time.
- * Many functions in R return lists.

Lists - Creating

```
family1 = list(husband="Fred", wife="Wilma", numofchildren=3, agesofkids=c(8,11,14))

length(family1) # Has 4 elements
[1] 4

family1
$husband
[1] "Fred"

$wife
[1] "Wilma"

$numofchildren
[1] 3

$agesofkids
[1] 8 11 14

str(family1)
List of 4
 $ husband      : chr "Fred"
 $ wife         : chr "Wilma"
 $ numofchildren: num 3
 $ agesofkids   : num [1:3] 8 11 14
```

Lists - Indexing

```
family1 = list(husband="Fred", wife="Wilma", numofchildren=3, agesofkids=c(8,11,14))

family1$agesofkids      # If the list elements have names then use "$" to access the element
[1] 8 11 14

family1$agesofkids[1:2]
[1] 8 11

sapply(family1,class)
      husband      wife numofchildren  agesofkids
"character" "character"    "numeric"    "numeric"

sapply(family1,length)
      husband      wife numofchildren  agesofkids
           1           1           1           3
```

If the list elements have no names then you have to use numeric indexing

```
family2 = list("Barney","Betty",2,c(4,6))

[[1]]
[1] "Barney"

[[2]]
[1] "Betty"

[[3]]
[1] 2

[[4]]
[1] 4 6
```

Lists - Indexing

If the list elements have no names then you have to use numeric indexing. But try to create lists with names as its easier to work with.

```
family2 = list("Barney","Betty",2,c(4,6))
```

```
family2[4]      # Accesses the 4th index and associated element  
[[1]]  
[1] 4 6
```

```
family2[[4]]    # Accesses the 4th element value only - more direct  
[1] 4 6
```

```
family2[3:4]    # Get 3rd and 4th indices and associate values  
[[1]]  
[1] 2
```

```
[[2]]  
[1] 4 6
```

Lists - Indexing

You can do "unlist" on any list to turn it into a vector. Since the list has mixed data types all of the elements of the vector will be converted to a single data type. In this case character

```
unlist(family1)
      husband      wife numofchildren  agesofkids1  agesofkids2
      "Fred"      "Wilma"          "3"          "8"          "11"
      agesofkids3
      "14"
```

```
as.numeric(unlist(family1))
[1] NA NA  3  8 11 14
```

Normally we don't create lists as a "standalone" object except in two major cases:

- 1) We are writing a function that does some interesting stuff and we want to return to the user a structure that has lots of information.
- 2) As a precursor to creating a data frame which is a hybrid between a list and a matrix. We'll investigate this momentarily.

Lists - Functions

R has lots of statistical functions that return lists of information. In fact this is the norm.

```
data(mtcars) # Load mtcars into the environment
```

```
mylm = lm(mpg ~ wt, data = mtcars)
```

```
print(mylm)
```

Call:

```
lm(formula = mpg ~ wt, data = mtcars)
```

Coefficients:

| | |
|-------------|--------|
| (Intercept) | wt |
| 37.285 | -5.344 |

```
# But there is a lot more information
```

```
typeof(mylm)
```

```
[1] "list"
```


Lists - Functions

R has lots of statistical functions that return lists of information. In fact this is the norm.

```
str(mylm,give.attr=F) # Lots of stuff here
```

List of 12

```
$ coefficients : Named num [1:2] 37.29 -5.34
$ residuals    : Named num [1:32] -2.28 -0.92 -2.09 1.3 -0.2 ...
$ effects      : Named num [1:32] -113.65 -29.116 -1.661 1.631 0.111 ...
$ rank         : int 2
$ fitted.values: Named num [1:32] 23.3 21.9 24.9 20.1 18.9 ...
$ assign       : int [1:2] 0 1
$ qr           :List of 5
..$ qr        : num [1:32, 1:2] -5.657 0.177 0.177 0.177 0.177 ...
..$ qraux      : num [1:2] 1.18 1.05
..$ pivot      : int [1:2] 1 2
..$ tol        : num 1e-07
..$ rank       : int 2
$ df.residual  : int 30
$ xlevels      : Named list()
$ call         : language lm(formula = mpg ~ wt, data = mtcars)
$ terms        :Classes 'terms', 'formula' length 3 mpg ~ wt
$ model        :'data.frame': 32 obs. of 2 variables:
..$ mpg        : num [1:32] 21 21 22.8 21.4 18.7 18.1 14.3 24.4 22.8 19.2 ...
..$ wt         : num [1:32] 2.62 2.88 2.32 3.21 3.44 ...
```

Lists - Functions

```
names(mylm)
[1] "coefficients" "residuals"    "effects"      "rank"
[5] "fitted.values" "assign"        "qr"           "df.residual"
[9] "xlevels"      "call"         "terms"        "model"

mylm$effects
(Intercept)          wt
-113.6497374 -29.1157217 -1.6613339  1.6313943  0.1111305 -0.3840041
-3.6072442  4.5003125  2.6905817  0.6111305 -0.7888695  1.1143917
 0.2316793 -1.6061571  1.3014525  2.2137818  6.0995633  7.3094734
 2.2421594  6.8956792 -2.2010595 -2.6694078 -3.4150859 -3.1915608
 2.7346556  0.8200064  0.5948771  1.7073457 -4.2045529 -2.4018616
-2.9072442 -0.6494289

# Some use the $ notation to extract desired information they want straight from the function call

lm(mpg ~ wt, data = mtcars)$coefficients
(Intercept)          wt
 37.285126    -5.344472
```

Lists - Functions

When we create our own functions we can package things up into a list and return things.

```
my.summary <- function(x) {  
  return.list = list()      # Declare the list  
  
  return.list$mean = mean(x)  
  return.list$sd = sd(x)  
  return.list$var = var(x)  
  
  return(return.list)  
}  
  
my.summary(1:10)  
  
$mean  
[1] 5.5  
  
$sd  
[1] 3.02765  
  
$var  
[1] 9.166667  
  
names(my.summary(1:10))  
[1] "mean" "sd"  "var"  
  
my.summary(1:10)$var      # Here we exploit the $ notation to get only what we want  
[1] 9.166667
```

Lists - Functions

Some other basic R functions will return a list - such as some of the character functions:

```
mystring = "This is a test"
```

```
mys = strsplit(mystring, " ")
```

```
str(mys)
```

```
List of 1
```

```
$ : chr [1:4] "This" "is" "a" "test"
```

```
mys
```

```
[[1]]
```

```
[1] "This" "is"    "a"     "test"
```

```
mys[[1]][1]
```

```
[1] "This"
```

```
mys[[1]][1:2]
```

```
[1] "This" "is"
```

```
unlist(mys)
```

```
[1] "This" "is"    "a"     "test"
```

Lists - Twitter

```
delta.tweets = searchTwitter('@delta', n = 100) # Uses the add-on twitterR package

class(delta.tweets)
[1] "list"

delta.tweets
[[1]]
[1] "sotsoy: Apparently if you use your frequent flier miles on @delta they stick you at the back of the plane on every flight next to the bathroom"

[[2]]
[1] "ImTooNonFiction: My @Delta flight has been delayed for the last 2 hrs. We've been on plane at gate for 2+ hours and no mention of a voucher or compensation"

[[3]]
[1] "ShaneNHara: @Delta and @DeltaAssist, thank you for a swift boarding process here at SEA en route to LAX. Taking care of your loyal flyers = appreciated."

[[4]]
[1] "NaiiOLLG: RT @TheRealNickMara: ThankYou @Delta for a great flight!! #Work!!!"

[[5]]
[1] "forbeslancaster: @bsideblog @Delta just saw a commercial highlighting delta awesome service. Totes NOT true"

..
```

Lists - Twitter

```
sapply(delta.tweets,function(x) x$getText()) # Pulls out the text of the tweet
```

```
[1] "Apparently if you use your frequent flier miles on @delta they stick you at the back of the plane on every flight next to the bathroom"
```

```
[2] "My @Delta flight has been delayed for the last 2 hrs. We've been on plane at gate for 2+ hours and no mention of a voucher or compensation"
```

```
[3] "@Delta and @DeltaAssist, thank you for a swift boarding process here at SEA en route to LAX. Taking care of your loyal flyers = appreciated."
```

```
[4] "RT @TheRealNickMara: ThankYou @Delta for a great flight!! #Work!!!"
```

```
[5] "@bsideblog @Delta just saw a commercial highlighting delta awesome service. Totes NOT true"
```

```
..  
..  
..
```

```
other results omitted due to obscenities...
```

Lists - sapply

Lastly, While we could use sapply to apply some statistical function across all elements of a list it might not make sense since you have different data types:

```
sapply(family1,mean)
      husband      wife numofchildren  agesofkids
      NA        NA           3           11
```

Warning messages:

```
1: In mean.default(X[[1L]], ...) :
  argument is not numeric or logical: returning NA
2: In mean.default(X[[2L]], ...) :
  argument is not numeric or logical: returning NA
```

We could write our own function to ignore non-numeric data:

```
sapply(family1, function(x) { if (is.numeric(x)) print(mean(x))})
[1] 3
[1] 11
$husband
NULL
$wife
NULL

$numofchildren
[1] 3

$agesofkids
[1] 11
```

Factors - Intro

R supports factors, which are a special data type for, among other things, managing categories of data.

“One of the most important uses of factors is in statistical modeling; since categorical variables enter into statistical models differently than continuous variables, storing data as factors insures that the modeling functions will treat such data correctly”.

Identifying categorical variables is usually straightforward. These are the variables by which you might want to summarize some continuous data.

Categorical variables usually take on a definite number of values.

Factors - Intro

Let's say we have some automobile data that tells us if a car has an automatic transmission (0) or a manual transmission (1). We store this into a vector called transvec

```
transvec = c(1,1,1,0,0,0,0,0,0,0,0,0,0,0,0,0,0,1,1,1,0,0,0,0,0,1,1,1,1,1,1,1)
```

```
table(transvec)    # Count 'em up. Which are Auto and Manual ?
```

```
transvec
```

```
 0  1
```

```
19 13
```

```
mytransfac = factor(transvec, levels = c(0,1), labels = c("Auto","Man"))
```

```
table(mytransfac)
```

```
mytransfac
```

```
Auto  Man
```

```
 19   13
```

```
levels(mytransfac)
```

```
[1] "Auto" "Man"
```

```
mytransfac
```

```
[1] Man  Man  Man  Auto Auto Auto Auto Auto Auto Auto Auto Auto Auto Auto Auto Auto
```

```
[16] Auto Auto Man  Man  Man  Auto Auto Auto Auto Auto Man  Man  Man  Man  Man
```

```
[31] Man  Man
```

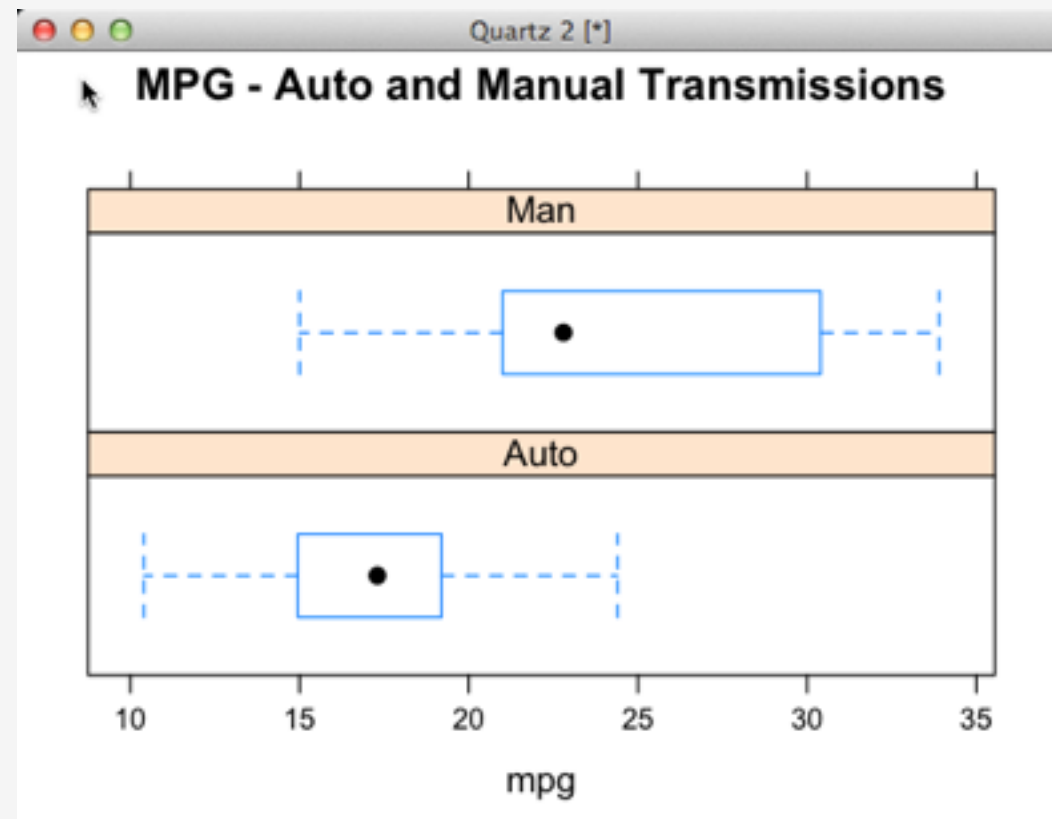

```
Levels: Auto Man
```

Factors - Intro

R knows how to handle factors when doing plots. Here we get an X/Y plot and a Box Plot with very little work since R knows that `mytransfac` is a factor

```
library(lattice)
xyplot(mpg~wt | mytransfac, mtcars, main="MPG vs Weight - Auto and Manual Transmissions")

bwplot(~mpg|mytransfac, mtcars, main="MPG - Auto and Manual Transmissions", layout=c(1,2))
```



Factors - Aggregation Preview

With our knowledge of factors and vectors we can do some basic aggregation using the `tapply` command.

We have a factor vector called `mytransfac`. Let's summarize some MPG data that corresponds to the automobiles used in the `mytransfac` vector. So for each car we have its MPG figure and whether it has an automatic or manual transmission.

```
mympg = c(21,21,22.8,21.4,18.7,18.1,14.3,24.4,22.8,19.2,17.8,16.4,17.3,15.2,10.4,  
          10.4,14.7,32.4,30.4,33.9,21.5,15.5,15.2,13.3,19.2,27.3,26,30.4,15.8,19.7,15,21.4)
```

```
tapply( continuous_value_to_summarize, factor_or_grouping_variable, function_for_summary)
```

```
tapply(mympg,mytransfac,mean)
```

```
      Auto      Man  
17.14737 24.39231
```

Factors - cut

It is sometimes useful to take a continuous variable and chop it up into intervals or categories for purposes of summary or grouping. R has a function to do this called “cut” to accomplish this. Let’s work through some examples to understand what is going on:

Let’s cut up the numbers between 1 and 10 into 4 intervals. It looks kind of messy:

```
cut(0:10,breaks=4)
```

```
[1] (-0.01,2.5] (-0.01,2.5] (-0.01,2.5] (2.5,5]      (2.5,5]      (2.5,5]      (5,7.5]      (5,7.5]
[7.5,10]      (7.5,10]
[11] (7.5,10]
Levels: (-0.01,2.5] (2.5,5] (5,7.5] (7.5,10]
```

```
table(cut(0:10,breaks=4))
```

```
(-0.01,2.5]      (2.5,5]      (5,7.5]      (7.5,10]
           3           3           2           3
```

Factors - cut

Well that was cool but people like to read labels:

```
my.cut = cut(0:10,breaks=4,labels=c("Q1","Q2","Q3","Q4"))  
[1] Q1 Q1 Q1 Q2 Q2 Q2 Q3 Q3 Q4 Q4 Q4  
Levels: Q1 Q2 Q3 Q4
```

```
table(my.cut)  
my.cut  
Q1 Q2 Q3 Q4  
 3  3  2  3
```

Factors - cut

But you can take finer-grained control over how the intervals are made.

```
quantile(0:10)
 0%  25%  50%  75% 100%
0.0  2.5  5.0  7.5 10.0

table(cut(0:10,breaks=quantile(0:10),include.lowest=TRUE))

 [0,2.5]  (2.5,5]  (5,7.5] (7.5,10]
      3         3         2         3
```

Factors - cut

Another example. Let's say we have some exam scores. Let's summarize them according to the typical US grading system. F: < 60, D: 60-70: C: 70-80: B:80-90 A:90-100

```
set.seed(123)
exam.score = runif(25,50,100)
```

```
cut(exam.score,breaks=c(50,60,70,80,90,100))
[1] (60,70] (80,90] (70,80] (90,100] (90,100] (50,60] (70,80] (90,100]
[9] (70,80] (70,80] (90,100] (70,80] (80,90] (70,80] (50,60] (90,100]
[17] (60,70] (50,60] (60,70] (90,100] (90,100] (80,90] (80,90] (90,100]
[25] (80,90]
Levels: (50,60] (60,70] (70,80] (80,90] (90,100]
```


```
cut(exam.score,breaks=c(50,60,70,80,90,100),labels=c("F","D","C","B","A"))
[1] D B C A A F C A C C A C B C F A D F D A A B B A B
Levels: F D C B A
```

```
my.table = table(cut(exam.score,breaks=c(50,60,70,80,90,100),labels=c("F","D","C","B","A")))
```

```
F D C B A
3 3 6 5 8
```

```
barchart(my.table,main="Grade BarChart",col=terrain.colors(5))
```

Factors - cut



Factors - cut

We have a small problem in that the intervals don't exactly match the grading scheme. In this scheme someone getting a grade of 90 will get a B although we intend for them to get an A. This is where you should be paying attention to the (and] characters. To make the interval exclude the "right side" of the interval we specify the "right=F" argument.

```
cut(exam.score,breaks=c(50,60,70,80,90,100))
[1] (60,70] (80,90] (70,80] (90,100] (90,100] (50,60] (70,80] (90,100]
[9] (70,80] (70,80] (90,100] (70,80] (80,90] (70,80] (50,60] (90,100]
[17] (60,70] (50,60] (60,70] (90,100] (90,100] (80,90] (80,90] (90,100]
[25] (80,90]
Levels: (50,60] (60,70] (70,80] (80,90] (90,100]
```

```
cut(exam.score,breaks=c(50,60,70,80,90,100),right=F)
[1] [60,70) [80,90) [70,80) [90,100) [90,100) [50,60) [70,80) [90,100)
[9] [70,80) [70,80) [90,100) [70,80) [80,90) [70,80) [50,60) [90,100)
[17] [60,70) [50,60) [60,70) [90,100) [90,100) [80,90) [80,90) [90,100)
[25] [80,90)
Levels: [50,60) [60,70) [70,80) [80,90) [90,100)
```

Factors - cut

So if you don't think that the cut command doesn't do something interesting then here is how you would have had to the last example with the exams:

```
exam.score = runif(25,50,100)

acount = 0
bcount = 0
ccount = 0
dcount = 0
fcount = 0
exam.score = runif(25,50,100)
for (ii in 1:length(exam.score)) {

  if (exam.score[ii] < 60) {fcount = fcount + 1} else
    if ((exam.score[ii] >= 60) & (exam.score[ii] < 70)) {dcount = dcount + 1} else
      if ((exam.score[ii] >= 70) & (exam.score[ii] < 80)) {ccount = ccount +1} else
        if ((exam.score[ii] >= 80) & (exam.score[ii] < 90)) {bcount = bcount +1} else
          if ((exam.score[ii] >= 90) & (exam.score[ii] <= 100)) {acount = acount +1}
}
cat("acount bcount ccount dcount fcount")
cat(acount,bcount,ccount,dcount,fcount)
acount bcount ccount dcount fcount
8 5 7 3 2
```

Factors - Ordered

Sometimes we want our factors to be ordered. For example we intuitively know that January comes before February and so on. Can we get R to create ordered factors ?

```
mons =c("Jan","Feb","Mar","Apr","May","Jun","Jan","Feb","May","Jun", "Apr","Mar")
```

```
my.fact.mons = factor(mons)
```

```
[1] Jan Feb Mar Apr May Jun Jan Feb May Jun Apr Mar  
Levels: Apr Feb Jan Jun Mar May
```

```
my.fact.mons[1] < my.fact.mons[2]
```

Warning message:

```
In Ops.factor(my.fact.mons[1], my.fact.mons[2]) :  
  < not meaningful for factors
```

```
levels(my.fact.mons)
```

```
[1] "Apr" "Feb" "Jan" "Jun" "Mar" "May"
```

<http://www.stat.berkeley.edu/classes/s133/factors>

Factors - Ordered

```
my.fact.mons = factor(mons, labels=c("Jan","Feb","Mar","Apr","May","Jun"),ordered=TRUE)
```

```
my.fact.mons
[1] Mar Feb May Jan Jun Apr Mar Feb Jun Apr Jan May
Levels: Jan < Feb < Mar < Apr < May < Jun
```

```
my.fact.mons[1] < my.fact.mons[2]
[1] FALSE
```

```
table(my.fact.mons)
my.fact.mons
Jan Feb Mar Apr May Jun
  2   2   2   2   2   2
```

```
levels(my.fact.mons)           # This is what we want !
[1] "Jan" "Feb" "Mar" "Apr" "May" "Jun"
```

<http://www.stat.berkeley.edu/classes/s133/factors>

Supplemental Factors - AOV example

Let's do an AOV on the mtcars data set variables MPG and number of gears the latter of which takes on the values 3,4,5. So it is well suited to be a factor.

```
mtcars$gear = factor(mtcars$gear) # Turn gear into a factor
aov.ex1 = aov(mpg ~ gear,mtcars)
summary(aov.ex1)
```

| | Df | Sum Sq | Mean Sq | F value | Pr(>F) |
|--------------|----|--------|---------|---------|---------------|
| factor(gear) | 2 | 483.24 | 241.622 | 10.901 | 0.0002948 *** |
| Residuals | 29 | 642.80 | 22.166 | | |

```
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
> print(model.tables(aov.ex1,"means"))
Tables of means
Grand mean

20.09062

gear
      3      4      5
16.11 24.53 21.38
rep 15.00 12.00  5.00

par(mfrow=c(2,2))
plot(aov.ex1)
```

Supplemental Factors - AOV example

Let's do an AOV on the mtcars data set variables MPG and number of gears the latter of which takes on the values 3,4,5. So it is well suited to be a factor.

```
my.tukey = TukeyHSD(aov.ex1,"gear") # Tukey Multiple Comparisons
```

```
my.tukey
```

```
Tukey multiple comparisons of means  
95% family-wise confidence level
```

```
Fit: aov(formula = mpg ~ gear, data = mtcars)
```

```
$gear
```

| | diff | lwr | upr | p adj |
|-----|-----------|------------|-----------|-----------|
| 4-3 | 8.426667 | 3.9234704 | 12.929863 | 0.0002088 |
| 5-3 | 5.273333 | -0.7309284 | 11.277595 | 0.0937176 |
| 5-4 | -3.153333 | -9.3423846 | 3.035718 | 0.4295874 |


Differences between Gears are significant at 5% level if the confidence interval around the estimation of the difference does not contain zero

```
plot(my.tukey)
```

Supplemental Factors - AOV example

Using R – Factors – Tukey–MultiComp Plot

95% family-wise confidence level



Dataframes



Dataframes

A **data frame** is a special type of list that contains data in a format that allows for easier manipulation, reshaping, and open-ended analysis.

Data frames are tightly coupled collections of variables. It is one of the more important constructs you will encounter when using R so learn all you can about it.

A data frame is an analogue to the Excel spreadsheet. In general this is the most popular construct for storing, manipulating, and analyzing data.

Data frames can be constructed from existing vectors, lists, or matrices. Many times they are created by reading in comma delimited files, (CSV files), using the `read.table` command.

Once you become accustomed to working with data frames, R becomes so much easier to use.

Dataframes

Here we have 4 vectors two of which are character and two of which are numeric.

We could work with them in the following fashion if we wanted to do some type of summary on them.

```
names = c("P1","P2","P3","P4","P5")
temp = c(98.2,101.3,97.2,100.2,98.5)
pulse = c(66,72,83,85,90)
gender = c("M","F","M","M","F")
```

We could write a for loop to get information for each patient but this isn't # so convenient or scalable.

```
for (ii in 1:length(gender)) {
  print.string = c(names[ii],temp[ii],pulse[ii],gender[ii])
  print(print.string)
}
```

```
[1] "P1"    "98.2" "66"    "M"
[1] "P2"    "101.3" "72"    "F"
[1] "P3"    "97.2" "83"    "M"
[1] "P4"    "100.2" "85"    "M"
[1] "P5"    "98.5" "90"    "F"
```

Dataframes

A data frame can be regarded as a matrix with columns possibly of differing modes and attributes. It may be displayed in matrix form, and its rows and columns extracted using matrix indexing conventions. Let's create a data frame:

```
names=c("P1","P2","P3","P4","P5")
temp=c(98.2,101.3,97.2,100.2,98.5)
pulse=c(66,72,83,85,90)
gender=c("M","F","M","M","F")
```

```
my_df = data.frame(names,temp,pulse,gender) # Much more flexible
my_df
```

| | names | temp | pulse | gender |
|---|-------|-------|-------|--------|
| 1 | P1 | 98.2 | 66 | M |
| 2 | P2 | 101.3 | 72 | F |
| 3 | P3 | 97.2 | 83 | M |
| 4 | P4 | 100.2 | 85 | M |
| 5 | P5 | 98.5 | 90 | F |

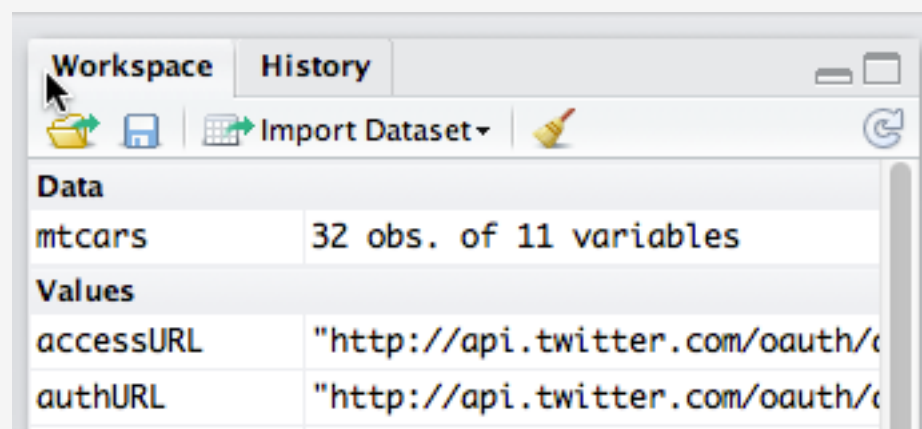
```
plot(my_df$pulse ~ my_df$temp,main="Pulse Rate",xlab="Patient",ylab="BPM")
```

```
mean(my_df[,2:3])
temp pulse
99.08 79.20
```

Dataframes

Once you have the data frame you could edit it with a GUI editor. Or you can use the Workspace Viewer/Editor in RStudio

```
data(mtcars) # This will load a copy of mtcars into your workspace.
```



The screenshot shows the RStudio Editor pane with the 'mtcars' dataset loaded. The title bar indicates 'mtcars' is the active file. The editor displays a table with 32 observations of 11 variables. The table has columns: row.names, mpg, cyl, disp, hp, drat, wt, qsec, vs, am, gear, and carb. The first 17 rows are visible, showing car models and their specifications.

| | row.names | mpg | cyl | disp | hp | drat | wt | qsec | vs | am | gear | carb |
|----|---------------------|------|-----|-------|-----|------|-------|-------|----|----|------|------|
| 1 | Mazda RX4 | 21.0 | 6 | 160.0 | 110 | 3.90 | 2.620 | 16.46 | 0 | 1 | 4 | 4 |
| 2 | Mazda RX4 Wag | 21.0 | 6 | 160.0 | 110 | 3.90 | 2.875 | 17.02 | 0 | 1 | 4 | 4 |
| 3 | Datsun 710 | 22.8 | 4 | 108.0 | 93 | 3.85 | 2.320 | 18.61 | 1 | 1 | 4 | 1 |
| 4 | Hornet 4 Drive | 21.4 | 6 | 258.0 | 110 | 3.08 | 3.215 | 19.44 | 1 | 0 | 3 | 1 |
| 5 | Hornet Sportabout | 18.7 | 8 | 360.0 | 175 | 3.15 | 3.440 | 17.02 | 0 | 0 | 3 | 2 |
| 6 | Valiant | 18.1 | 6 | 225.0 | 105 | 2.76 | 3.460 | 20.22 | 1 | 0 | 3 | 1 |
| 7 | Duster 360 | 14.3 | 8 | 360.0 | 245 | 3.21 | 3.570 | 15.84 | 0 | 0 | 3 | 4 |
| 8 | Merc 240D | 24.4 | 4 | 146.7 | 62 | 3.69 | 3.190 | 20.00 | 1 | 0 | 4 | 2 |
| 9 | Merc 230 | 22.8 | 4 | 140.8 | 95 | 3.92 | 3.150 | 22.90 | 1 | 0 | 4 | 2 |
| 10 | Merc 280 | 19.2 | 6 | 167.6 | 123 | 3.92 | 3.440 | 18.30 | 1 | 0 | 4 | 4 |
| 11 | Merc 280C | 17.8 | 6 | 167.6 | 123 | 3.92 | 3.440 | 18.90 | 1 | 0 | 4 | 4 |
| 12 | Merc 450SE | 16.4 | 8 | 275.8 | 180 | 3.07 | 4.070 | 17.40 | 0 | 0 | 3 | 3 |
| 13 | Merc 450SL | 17.3 | 8 | 275.8 | 180 | 3.07 | 3.730 | 17.60 | 0 | 0 | 3 | 3 |
| 14 | Merc 450SLC | 15.2 | 8 | 275.8 | 180 | 3.07 | 3.780 | 18.00 | 0 | 0 | 3 | 3 |
| 15 | Cadillac Fleetwood | 10.4 | 8 | 472.0 | 205 | 2.93 | 5.250 | 17.98 | 0 | 0 | 3 | 4 |
| 16 | Lincoln Continental | 10.4 | 8 | 460.0 | 215 | 3.00 | 5.424 | 17.82 | 0 | 0 | 3 | 4 |
| 17 | Chrysler Imperial | 14.7 | 8 | 440.0 | 230 | 3.23 | 5.345 | 17.42 | 0 | 0 | 3 | 4 |

Dataframes

R comes with a variety of built-in data sets that are very useful for getting used to data sets and how to manipulate them.

```
library(help="datasets")
```

```
# Gives detailed descriptions on available data sets
```

| | |
|----------------|---|
| AirPassengers | Monthly Airline Passenger Numbers 1949-1960 |
| BJsales | Sales Data with Leading Indicator |
| BOD | Biochemical Oxygen Demand |
| CO2 | Carbon Dioxide Uptake in Grass Plants |
| ChickWeight | Weight versus age of chicks on different diets |
| DNase | Elisa assay of DNase |
| EuStockMarkets | Daily Closing Prices of Major European Stock Indices, 1991-1998 |
| Formaldehyde | Determination of Formaldehyde |
| HairEyeColor | Hair and Eye Color of Statistics Students |

```
help(mtcars)    # Get details on a given data set
```

Dataframes

Let's focus on one of the built in sets. Its called "mtcars". The data was extracted from the 1974 _Motor Trend_ US magazine, and comprises fuel consumption and 11 aspects of automobile design and performance for 32 automobiles (1973-74 models).

```
[, 1] mpg    Miles/(US) gallon
[, 2] cyl    Number of cylinders
[, 3] disp   Displacement (cu.in.)
[, 4] hp     Gross horsepower
[, 5] drat   Rear axle ratio
[, 6] wt     Weight (lb/1000)
[, 7] qsec   1/4 mile time
[, 8] vs     V/S
[, 9] am     Transmission (0 = automatic, 1 = manual)
[,10] gear   Number of forward gears
[,11] carb   Number of carburetors
```

One way to get the type of each column

```
> sapply(mtcars, class)
      mpg      cyl      disp      hp      drat      wt      qsec      vs
"numeric" "numeric" "numeric" "numeric" "numeric" "numeric" "numeric" "numeric"
      am      gear      carb
"numeric" "numeric" "numeric"
```

Dataframes

The data was extracted from the 1974 _Motor Trend_ US magazine, and comprises fuel consumption and 11 aspects of automobile design and performance for 32 automobiles (1973-74 models).

```
str(mtcars)
'data.frame':   32 obs. of  11 variables:
 $ mpg : num  21 21 22.8 21.4 18.7 18.1 14.3 24.4 22.8 19.2 ...
 $ cyl : num  6 6 4 6 8 6 8 4 4 6 ...
 $ disp: num  160 160 108 258 360 ...
 $ hp  : num  110 110 93 110 175 105 245 62 95 123 ...
 $ drat: num  3.9 3.9 3.85 3.08 3.15 2.76 3.21 3.69 3.92 3.92 ...
 $ wt  : num  2.62 2.88 2.32 3.21 3.44 ...
 $ qsec: num  16.5 17 18.6 19.4 17 ...
 $ vs  : num  0 0 1 1 0 1 0 1 1 1 ...
 $ am  : num  1 1 1 0 0 0 0 0 0 0 ...
 $ gear: num  4 4 4 3 3 3 3 4 4 4 ...
 $ carb: num  4 4 1 1 2 1 4 2 2 4 ...
```

```
nrow(mtcars)      # How many rows does it have ?
[1] 32
```

```
ncol(mtcars)      # How many columns are there ?
[1] 11
```

Dataframes

```
rownames(mtcars)
[1] "Mazda RX4"           "Mazda RX4 Wag"       "Datsun 710"
[4] "Hornet 4 Drive"      "Hornet Sportabout"   "Valiant"
..
[19] "Honda Civic"         "Toyota Corolla"      "Toyota Corona"
[22] "Dodge Challenger"    "AMC Javelin"         "Camaro Z28"
[25] "Pontiac Firebird"    "Fiat X1-9"           "Porsche 914-2"
[28] "Lotus Europa"        "Ford Pantera L"      "Ferrari Dino"
[31] "Maserati Bora"       "Volvo 142E"
```

```
rownames(mtcars) = 1:32
```

```
head(mtcars)
  mpg  cyl  disp  hp drat   wt  qsec vs transmission gear carb
1 21.0   6  160  110 3.90 2.62 16.5  0             1     4     4
2 21.0   6  160  110 3.90 2.88 17.0  0             1     4     4
```

```
rownames(mtcars) = paste("car",1:32,sep="_")
```

```
head(mtcars)
  mpg  cyl  disp  hp drat   wt  qsec vs transmission gear carb
car_1 21.0   6  160  110 3.90 2.62 16.5  0             1     4     4
car_2 21.0   6  160  110 3.90 2.88 17.0  0             1     4     4
car_3 22.8   4  108   93 3.85 2.32 18.6  1             1     4     1
```


Dataframes

There are various ways to **select, remove, or exclude** rows and columns from a data frame.

```
mtcars[, -11]
```

| | mpg | cyl | disp | hp | drat | wt | qsec | vs | am | gear |
|---------------|------|-----|------|-----|------|-------|-------|----|----|------|
| Mazda RX4 | 21.0 | 6 | 160 | 110 | 3.90 | 2.620 | 16.46 | 0 | 1 | 4 |
| Mazda RX4 Wag | 21.0 | 6 | 160 | 110 | 3.90 | 2.875 | 17.02 | 0 | 1 | 4 |
| Datsun 710 | 22.8 | 4 | 108 | 93 | 3.85 | 2.320 | 18.61 | 1 | 1 | 4 |

```
mtcars      # Notice that carb is included
```

| | mpg | cyl | disp | hp | drat | wt | qsec | vs | am | gear | carb |
|---------------|------|-----|-------|-----|------|-------|-------|----|----|------|------|
| Mazda RX4 | 21.0 | 6 | 160.0 | 110 | 3.90 | 2.620 | 16.46 | 0 | 1 | 4 | 4 |
| Mazda RX4 Wag | 21.0 | 6 | 160.0 | 110 | 3.90 | 2.875 | 17.02 | 0 | 1 | 4 | 4 |
| Datsun 710 | 22.8 | 4 | 108.0 | 93 | 3.85 | 2.320 | 18.61 | 1 | 1 | 4 | 1 |

```
mtcars[, -3:-5] # Print all columns except for columns 3 through 5
```

| | mpg | cyl | wt | qsec | vs | am | gear | carb |
|---------------|------|-----|-------|-------|----|----|------|-----------|
| Mazda RX4 | 21.0 | 6 | 2.620 | 16.46 | 0 | 1 | 4 | 0.6020600 |
| Mazda RX4 Wag | 21.0 | 6 | 2.875 | 17.02 | 0 | 1 | 4 | 0.6020600 |
| Datsun 710 | 22.8 | 4 | 2.320 | 18.61 | 1 | 1 | 4 | 0.0000000 |

```
mtcars[, c(-3, -5)] # Print all columns except for columns 3 AND 5
```

| | mpg | cyl | hp | wt | qsec | vs | am | gear | carb |
|---------------|------|-----|-----|-------|-------|----|----|------|-----------|
| Mazda RX4 | 21.0 | 6 | 110 | 2.620 | 16.46 | 0 | 1 | 4 | 0.6020600 |
| Mazda RX4 Wag | 21.0 | 6 | 110 | 2.875 | 17.02 | 0 | 1 | 4 | 0.6020600 |
| Datsun 710 | 22.8 | 4 | 93 | 2.320 | 18.61 | 1 | 1 | 4 | 0.0000000 |

Dataframes

There are various ways to **select, remove, or exclude** rows and columns from a data frame.

```
mtcars[mtcars$mpg >= 30.0,]
```

| | mpg | cyl | disp | hp | drat | wt | qsec | vs | am | gear | carb |
|----------------|------|-----|------|-----|------|-------|-------|----|----|------|------|
| Fiat 128 | 32.4 | 4 | 78.7 | 66 | 4.08 | 2.200 | 19.47 | 1 | 1 | 4 | 1 |
| Honda Civic | 30.4 | 4 | 75.7 | 52 | 4.93 | 1.615 | 18.52 | 1 | 1 | 4 | 2 |
| Toyota Corolla | 33.9 | 4 | 71.1 | 65 | 4.22 | 1.835 | 19.90 | 1 | 1 | 4 | 1 |
| Lotus Europa | 30.4 | 4 | 95.1 | 113 | 3.77 | 1.513 | 16.90 | 1 | 1 | 5 | 2 |

```
mtcars[mtcars$mpg >= 30.0,2:6]
```

| | mpg | cyl | disp | hp | drat |
|----------------|------|-----|------|-----|------|
| Fiat 128 | 32.4 | 4 | 78.7 | 66 | 4.08 |
| Honda Civic | 30.4 | 4 | 75.7 | 52 | 4.93 |
| Toyota Corolla | 33.9 | 4 | 71.1 | 65 | 4.22 |
| Lotus Europa | 30.4 | 4 | 95.1 | 113 | 3.77 |

```
mtcars[mtcars$mpg >= 30.0 & mtcars$cyl < 6,]
```

| | mpg | cyl | disp | hp | drat | wt | qsec | vs | am | gear | carb |
|----------------|------|-----|------|-----|------|-------|-------|----|----|------|------|
| Fiat 128 | 32.4 | 4 | 78.7 | 66 | 4.08 | 2.200 | 19.47 | 1 | 1 | 4 | 1 |
| Honda Civic | 30.4 | 4 | 75.7 | 52 | 4.93 | 1.615 | 18.52 | 1 | 1 | 4 | 2 |
| Toyota Corolla | 33.9 | 4 | 71.1 | 65 | 4.22 | 1.835 | 19.90 | 1 | 1 | 4 | 1 |
| Lotus Europa | 30.4 | 4 | 95.1 | 113 | 3.77 | 1.513 | 16.90 | 1 | 1 | 5 | 2 |

Dataframes

Find all rows that correspond to Automatic and Count them

```
mtcars[mtcars$am==0,]
      mpg  cyl  disp  hp drat   wt  qsec vs  am  gear  carb
Hornet 4 Drive 21.4   6 258.0 110 3.08 3.215 19.44 1   0     3     1
Hornet Sportabout 18.7   8 360.0 175 3.15 3.440 17.02 0   0     3     2
Valiant      18.1   6 225.0 105 2.76 3.460 20.22 1   0     3     1
Duster 360    14.3   8 360.0 245 3.21 3.570 15.84 0   0     3     4
Merc 240D     24.4   4 146.7  62 3.69 3.190 20.00 1   0     4     2
Merc 230     22.8   4 140.8  95 3.92 3.150 22.90 1   0     4     2
..
..

nrow(mtcars[mtcars$am == 0,])
[1] 19

nrow(mtcars[mtcars$am == 1,])
[1] 13
```

Dataframes

Many times data will be read in from a comma delimited ,("CSV"), file exported from Excel. The file can be read from local storage or from the Web.

```
url = "http://www.bimcore.emory.edu/BIOS560R/DATA.DIR/hsb2.csv"
```

```
data1 = read.table(url,header=T,sep=",")
```

```
head(data1)
```

| | gender | id | race | ses | schtyp | prgtype | read | write | math | science | socst |
|---|--------|-----|------|-----|--------|----------|------|-------|------|---------|-------|
| 1 | 0 | 70 | 4 | 1 | 1 | general | 57 | 52 | 41 | 47 | 57 |
| 2 | 1 | 121 | 4 | 2 | 1 | vocati | 68 | 59 | 53 | 63 | 61 |
| 3 | 0 | 86 | 4 | 3 | 1 | general | 44 | 33 | 54 | 58 | 31 |
| 4 | 0 | 141 | 4 | 3 | 1 | vocati | 63 | 44 | 47 | 53 | 56 |
| 5 | 0 | 172 | 4 | 2 | 1 | academic | 47 | 52 | 57 | 53 | 61 |
| 6 | 0 | 113 | 4 | 2 | 1 | academic | 44 | 52 | 51 | 63 | 61 |

```
sapply(data1,class)      # Applies the "Class" function to all columns
gender      id      race      ses      schtyp      prgtype      read      write
"integer" "integer" "integer" "integer" "integer"  "factor" "integer" "integer"

      math      science      socst
"integer" "integer" "integer"
```

Dataframes

Or you can use the “colClasses” argument when calling `read.table()` which allows you to set the variable type as you read in the data. It takes a bit of work up front but is worth it since it requires you to think about what you want/need your variable types. You can always change the types after the fact as in the previous example.

```
myclasses = c("character", "integer", "integer",  
              "integer", "character", "factor", "integer", "integer",  
              "integer", "integer", "numeric")  
  
data1 = read.table("http://www.bimcore.emory.edu/BIOS560R/DATA.DIR/hsb2.csv",  
                  header=T,  
                  sep=";",  
                  colClasses = myclasses)
```

```
sapply(data1, class)  
      gender      id      race      ses      schtyp      prgtype  
"character" "integer" "integer" "integer" "character" "factor"  
  
      read      write      math      science      socst  
"integer" "integer" "integer" "integer" "numeric"
```

Dataframes

Back to the mtcars data frame. What columns appear to be candidates for a factor ?

It would be variables who have only "a few" number of different values. If we do something like this we can get an idea. Looks like the last 4 columns might be what they want.

```
str(mtcars)
'data.frame': 32 obs. of 11 variables:
 $ mpg : num 21 21 22.8 21.4 18.7 18.1 14.3 24.4 22.8 19.2 ...
 $ cyl : num 6 6 4 6 8 6 8 4 4 6 ...
 $ disp: num 160 160 108 258 360 ...
 $ hp : num 110 110 93 110 175 105 245 62 95 123 ...
 $ drat: num 3.9 3.9 3.85 3.08 3.15 2.76 3.21 3.69 3.92 3.92 ...
 $ wt : num 2.62 2.88 2.32 3.21 3.44 ...
 $ qsec: num 16.5 17 18.6 19.4 17 ...
 $ vs : num 0 0 1 1 0 1 0 1 1 1 ...
 $ am : num 1 1 1 0 0 0 0 0 0 0 ...
 $ gear: num 4 4 4 3 3 3 3 4 4 4 ...
 $ carb: num 4 4 1 1 2 1 4 2 2 4 ...
```

```
unique(mtcars$am) # Tells us what the unique values are
[1] 1 0
```

Dataframes

Back to the mtcars data frame. What columns appear to be candidates for a factor ?

It would be variables who have only "a few" number of different values. If we do something like this we can get an idea. Looks like the last 4 columns might be what they want.

```
sapply(mtcars[,8:11], unique) # applies the unique function to columns 8-11 inclusive
```

```
$vs
```

```
[1] 0 1
```

```
$am
```

```
[1] 1 0
```

```
$gear
```

```
[1] 4 3 5
```

```
$carb
```

```
[1] 4 1 2 3 6 8
```

Dataframes

```
mtcars$am = factor(mtcars$am, levels = c(0,1), labels = c("Auto","Man"))
```

```
str(mtcars$am)
```

```
Factor w/ 2 levels "Auto","Man": 2 2 2 1 1 1 1 1 1 1 ...
```

```
# See what we have now !
```

```
head(mtcars,5)
```

| | mpg | cyl | disp | hp | drat | wt | qsec | vs | am | gear | carb |
|-------------------|------|-----|------|-----|------|-------|-------|----|------|------|------|
| Mazda RX4 | 21.0 | 6 | 160 | 110 | 3.90 | 2.620 | 16.46 | 0 | Man | 4 | 4 |
| Mazda RX4 Wag | 21.0 | 6 | 160 | 110 | 3.90 | 2.875 | 17.02 | 0 | Man | 4 | 4 |
| Datsun 710 | 22.8 | 4 | 108 | 93 | 3.85 | 2.320 | 18.61 | 1 | Man | 4 | 1 |
| Hornet 4 Drive | 21.4 | 6 | 258 | 110 | 3.08 | 3.215 | 19.44 | 1 | Auto | 3 | 1 |
| Hornet Sportabout | 18.7 | 8 | 360 | 175 | 3.15 | 3.440 | 17.02 | 0 | Auto | 3 | 2 |

```
tapply(mtcars$mpg,mtcars$am,mean)
```

```
      Auto      Man  
17.14737 24.39231
```

```
tapply(mtcars$mpg,mtcars$am,quantile)
```

```
$Auto
```

```
  0%   25%   50%   75%  100%  
10.40 14.95 17.30 19.20 24.40
```

```
$Man
```

```
  0%   25%   50%   75%  100%  
15.0 21.0 22.8 30.4 33.9
```