
Introduction to R

Steve Pittard

Dept of Biostatistics and Bioinformatics

Emory University

wsp@emory.edu

Introduction - Motivations

Why R ? Occupational Outlook Handbook, 2011

"Employment in professional, scientific, and technical services is projected to grow by 34 percent, adding about 2.7 million new jobs by 2018. In particular Computer and mathematical science occupations are projected to add almost 785,700 new jobs from 2008 to 2018.

"As a group, these occupations are expected to grow more than twice as fast as the average for all occupations in the economy. Demand for workers in computer and mathematical occupations will be driven by the continuing need for businesses, government agencies, and other organizations to adopt and utilize the latest technologies".

Introduction - Motivations

GeneTech / Biostatistician, "Experience with statistical software packages such as **R**, **S-Plus**, SAS, and JMP. <http://jobs.gene.com/job/00376705/>"

ICF-CDC / Biostatistician "Specialized training in **S-Plus** ,(R's commercial antecedent) ,SAS, SPSS, BMDP, and/or other statistical analysis software"

XDx Molecular Diagnostics / Biostatistician Consultant - "Expert in **R**, or other statistical software packages"

IveyExec /Principal Statistician, "Knowledge of statistical software packages: SAS Splus/R, Matlab, C++, Visual Basic"

Tempero Pharmaceuticals / Genomics Data Analyst, Statistical programming using MATLAB and/or **R** "

Lockheed Martin / Bioinformatics Programmer - "Software application development includes using SAS or **S-Plus/R** and a working knowledge of relational database design"

Introduction - Motivations

Data Analysts Captivated by R's Power The New York Times



Stuart Iselt for The New York Times

R first appeared in 1996, when the statistics professors Robert Gentleman, left, and Ross Ihaka released the code as a free software package.

Introduction - Motivations

The New York Times

Business Computing

- "Companies as diverse as Google, Pfizer, Merck, Bank of America, and the InterContinental Hotels Group and Shell use R."
- R is really important to the point that it's hard to overvalue it," said Daryl Pregibon, a research scientist at Google, which uses the software widely. "It allows statisticians to do very intricate and complicated analyses without knowing the blood and guts of computing systems."
- R has really become the second language for people coming out of grad school now, and there's an amazing amount of code being written for it," said Max Kuhn, associate director of nonclinical statistics at Pfizer.
- Close to 1,600 different packages reside on just one of the many Web sites devoted to R, and the number of packages has grown exponentially.

Introduction - Motivations

- Facebook's Data Team used R to answer two questions about new users: (i) which data points predict whether a user will stay? and (ii) if they stay, which data points predict how active they'll be after three months?
- Facebook used recursive partitioning to infer that two data points are significantly predictive of whether a user remains on Facebook: (i) having more than one session as a new user, and (ii) entering basic profile information.
- For the second question, they fit data to a logistic model using a least angle regression approach to find that activity was predicted by (i) how often a user was reached out to by others, (ii) frequency of third party application use, and (iii) how forthcoming a user was on the site. <http://www.dataspora.com/2009/02/predictive-analytics-using-r/>

Introduction - Motivations

- Google determines the effectiveness of display ads for its customers. Using observational data from more than 10 million web users, Google compares the search behavior of people who were exposed to the display ad (i.e. those that never visited a web page displaying the ad) to similar users who did see the ad, to figure out how many additional people visit the advertiser's web site as a result of seeing the display ad.
- Conditional regression models are used to evaluate the factors that lead to user satisfaction of Google products, such as when users are surveyed on satisfaction with search reports, or when users are asked to rate YouTube videos.

<http://blog.revolutionanalytics.com/2011/08/google-r-effective-ads.html>

Introduction - Motivations

- R is an interactive framework for data and statistical analysis that also happens to have a programming language.
- Compare this to languages such as Perl, Python, and Java that have data analysis addons
- Which language to use ? Exploit the strengths of all of them but if data analysis is a big part of the work then consider using R as part of the "pipeline".
- R can be well integrated with Excel, existing C,C++, Perl, Python, XML, and FORTRAN programs. **See** www.omegahat.com
- Most of the effort in using R relates to shaping the data for analysis and understanding the available functions and packages - deep programming skills are **not** required but are definitely rewarded.

Introduction - Motivations

When talking about user friendliness of computer software I like the analogy of cars vs. busses: [...] Using this analogy programs like SPSS are busses, easy to use for the standard things, but very frustrating if you want to do something that is not already preprogrammed. R is a 4-wheel drive SUV (though environmentally friendly) with a bike on the back, a kayak on top, good walking and running shoes in the passenger seat, and mountain climbing and spelunking gear in the back. R can take you anywhere you want to go if you take time to learn how to use the equipment, but that is going to take longer than learning where the bus stops are in SPSS. – Greg Snow R-help (May 2006)

See Introduction to the R Statistical Computing Environment

<http://www.slideshare.net/izahn/rintro>

Introduction - Motivations

- Vast capabilities, wide range of statistical and graphical techniques
- Written primarily by statisticians
- FREE as in free beer: no cost
- FREE as in free speech: collaborative development
- Excellent community support: mailing list, blogs, tutorials
- Easy to extend by writing new functions

Taken from Introduction to the R Statistical Computing Environment

<http://www.slideshare.net/izahn/rintro>

Introduction - Motivations

- The R console
 - Displays command history and results
 - Commands can be typed directly in the console
 - R Console work disappears once session is closed
- A text editor
 - A plain text editor for writing R code
 - Good ones will have syntax highlighting, parentheses matching etc.
 - Anything that modifies your data should be done in a text editor
- Graphics windows
 - View, re-size, and save graphics
 - A good GUI will allow you to cycle through graph history
- Work-space viewer
 - Some GUIs have work-space browsers that allow you to see stored objects
 - Very helpful if you are absentminded like me and frequently forget what names you gave your data!

Taken from Introduction to the R Statistical Computing Environment

<http://www.slideshare.net/izahn/rintro>

Installing

For Windows:

- 1) Go to <http://www.r-project.org>
- 2) Click on "CRAN" in the left navigation bar. See the list of mirror sites setup by country
- 3) Pick a site in the US such as <http://cran.cnr.berkeley.edu/>
- 4) Click on "Download R for Windows"
- 5) Then click on the "base" link
- 6) Now click on the link at the top of page (e.g. Download R 2.13.1 for Windows)
- 7) After the download completes then double-click the .exe file to initiate the installation. Answer the questions as they are presented.
- 8) To start R, click on Start -> All Programs -> R
- 9) **OPTIONAL:** Consider installing R Studio. While its not a requirement it is a great GUI for R. See <http://rstudio.org/download/desktop>

Installing

For Apple OSX:

- 1) Go to <http://www.r-project.org>
- 2) Click on "CRAN" in the left navigation bar. See the list of mirror sites setup by country
- 3) Pick a site in the US such as <http://cran.cnr.berkeley.edu/>
- 4) Click on "Download R for Mac OSX". This will start the download
- 5) After the download completes then double-click the .pkg file to initiate the installation. Answer the questions as they are presented.
- 6) To start R, Click on the icon in the Applications dock. You can also launch a Terminal window and type "R" within it.
- 7) **OPTIONAL:** Consider installing R Studio. While its not a requirement it is a great GUI for R. See <http://rstudio.org/download/desktop>

Installing - Packages

It is important to note that R comes with a base set of packages as part of every install.

```
> getOption("defaultPackages")
[1] "datasets" "utils"      "grDevices" "graphics" "stats"      "methods"

> search()
[1] ".GlobalEnv" "package:lattice" "package:stats" "package:graphics"
"package:grDevices" "package:utils"
[7] "package:datasets" "package:methods" "Autoloads"      "package:base"
```

While this might not look like a lot, each package has lots of useful stuff contained therein. If you wanted to drill down into one of these, here is one way to do it:

```
> library(help="stats")
```

Installing - Packages

```
> library(help="stats")
      Information on package 'stats'
```

Description:

```
Package:      stats
Version:      2.15.1
Priority:      base
Title:        The R Stats Package
Author:       R Core Team and contributors worldwide
Maintainer:   R Core Team <R-core@r-project.org>
Description:  R statistical functions
License:      Part of R 2.15.1
Built:        R 2.15.1; universal-apple-darwin9.8.0; 2012-06-22 19:10:14 UTC; unix
```

Index:

<code>.checkMFClasses</code>	Functions to Check the Type of Variables passed to Model Frames
<code>AIC</code>	Akaike's An Information Criterion
<code>ARMAacf</code>	Compute Theoretical ACF for an ARMA Process
<code>ARMAtoMA</code>	Convert ARMA Process to Infinite MA Process
<code>Beta</code>	The Beta Distribution
<code>Binomial</code>	The Binomial Distribution
<code>Box.test</code>	Box-Pierce and Ljung-Box Tests

Installing - Packages

Many packages also come with test and example data sets that can be very useful when attempting to perfect your knowledge of the various functions. To see what test data sets are available in a give package, do something like:

```
> search()
[1] ".GlobalEnv"          "package:lattice"    "package:stats"
"package:graphics"    "package:grDevices" "package:utils"
[7] "package:datasets"    "package:methods"    "Autoloads"
"package:base"
```

```
> data(package="stats")
```

Data sets in package 'datasets':

AirPassengers

1949-1960

BJsales

BJsales.lead (BJsales)

BOD

CO2

DNase

EuStockMarkets

Stock Indice

..

Monthly Airline Passenger Numbers

Sales Data with Leading Indicator

Sales Data with Leading Indicator

Biochemical Oxygen Demand

Carbon Dioxide Uptake in Grass Plants

Elisa assay of DNase

Daily Closing Prices of Major European

Installing - CRAN

While the base R package has many cool functions, you will invariably want to install many of the add-on packages available from CRAN.

To obtain information on the wide variety of packages available for then visit <http://cran.cnr.berkeley.edu> Here are some of the available package categories. Use the “search” function to drill down:

Bayesian	Bayesian Inference
ChemPhys	Chemometrics and Computational Physics
ClinicalTrials	Clinical Trial Design, Monitoring, and Analysis
Cluster	Cluster Analysis & Finite Mixture Models
Distributions	Probability Distributions
Econometrics	Computational Econometrics
Environmetrics	Analysis of Ecological and Environmental Data
ExperimentalDesign	Design of Experiments (DoE) & Analysis of Experimental Data
Finance	Empirical Finance
Genetics	Statistical Genetics
Graphics	Graphic Displays & Dynamic Graphics & Graphic Devices & Visualization
gR	gRaphical Models in R
HighPerformanceComputing	High-Performance and Parallel Computing with R

Installing - CRAN

If you are using an installation of R on Windows you can use a GUI to help you identify and install packages. However, if you are using a straight terminal interface such as when on Linux or OSX Terminal then you will need to something like the following.

Let's say that we have decided that we need the capabilities offered by the "actuar" function on CRAN. This package relates to Actuarial functions. To install it, start R.

```
> install.packages("actuar",dependencies=TRUE)
--- Please select a CRAN mirror for use in this session ---
Loading Tcl/Tk interface ... done

trying URL 'http://mirrors.nics.utk.edu/cran/bin/macosx/leopard/contrib/2.15/actuar\_1.1-5.tgz'
Content type 'application/x-gzip' length 1837121 bytes (1.8 Mb)
opened URL
=====
downloaded 1.8 Mb

> library(actuar)    # Brings the package into the workspace
```


Installing - CRAN

If you are using an installation of R on Windows you can use a GUI to help you identify and install packages. However, if you are using a straight terminal interface such as when on Linux or OSX Terminal then you will need to something like the following.

```
> library(actuar)    # Brings the package into the workspace

> search()
[1] ".GlobalEnv"          "package:actuar"      "package:lattice"
"package:stats"        "package:graphics"   "package:grDevices"
[7] "package:utils"       "package:datasets"   "package:methods"
"Autoloads"            "package:base"
```

Installing - CRAN

On occasion, (not often), you will need to install a package from a specific repository. For example, the OmegaHat project at www.omeghat.org has many interesting add-on packages that aren't necessarily linked into CRAN. To install one of these packages you would do something like:

```
> install.packages("GeoIP", repos = "http://www.omeghat.org/R")
```

On Linux and OSX it is possible, (though unlikely), that you will have to install a package using the "tar/gzipped" version of the file. This might happen if a colleague has written a package that hasn't been registered on CRAN. To do the install you do something like:

```
$ R CMD INSTALL GeoIP.tar.gz
```

It is also possible that the above **might** trigger a permissions error in which case you will have to use the "sudo" command or get the administrator of your system to help.

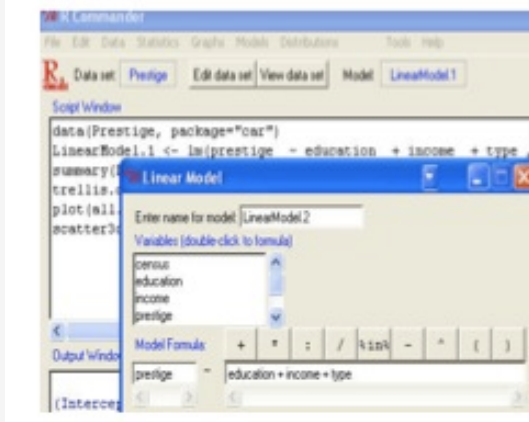
```
$ sudo R CMD INSTALL GeoIP.tar.gz
```

GUI ?

R is provided with a command line interface (CLI), which is the preferred user interface for power users because it allows direct control on calculations and is flexible. However, good knowledge of the language is required.

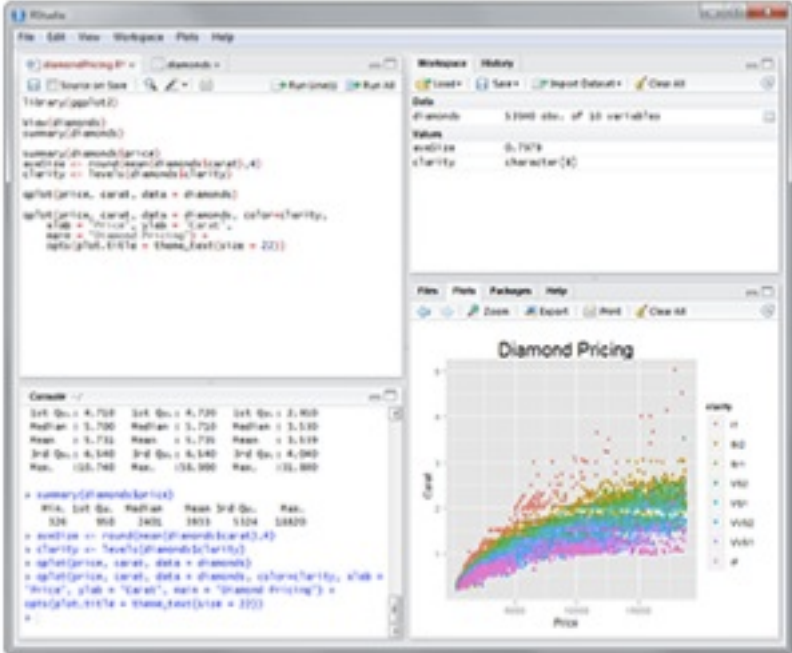
Because of this, the CLI can be intimidating for beginners. Thankfully, there are a number of R GUIs out there to make your life easier: R Commander, R-Studio, JGR, SciViews-R, R.app (OSX), Tinn-R (Windows only), ECLIPSE with STATET, EMACs

Note that on Windows you get default GUI that provides you with a reasonable amount of menu-drive capability. However, you could also use another GUI if you so desired.

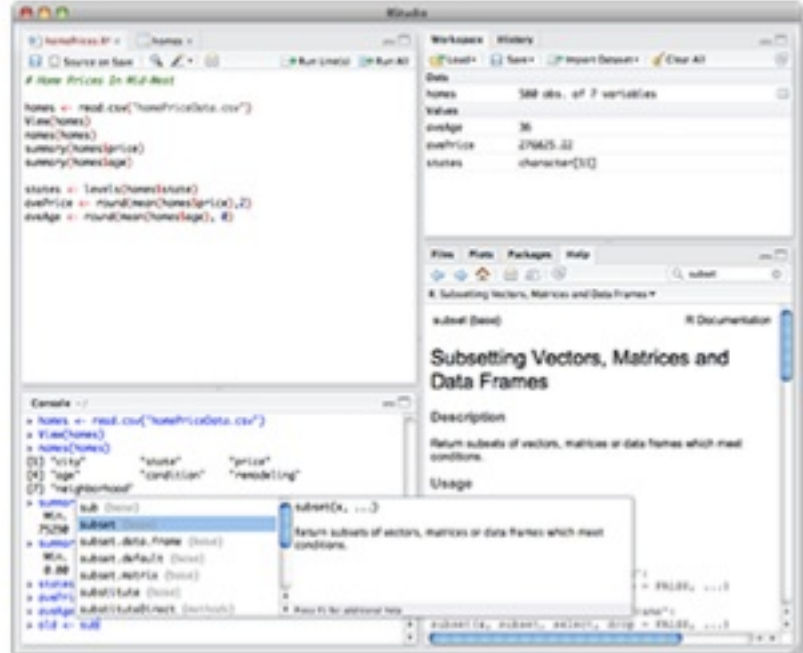


GUI ?

I like to use R-Studio. See <http://rstudio.org>



Windows



Mac OS X

GUI ?

I like to use R-Studio. See <http://rstudio.org>

Powerful productivity tools

- Syntax highlighting, code completion, and smart indentation
- Execute R code directly from the source editor
- Easily manage multiple working directories using projects
- Quickly navigate code using typeahead search and go to definition

An IDE built for R

- Workspace browser and data viewer
- Plot history, zooming, and flexible image and PDF export
- Integrated R help and documentation
- Sweave authoring including one-click PDF preview
- Searchable command history

Open and compatible

- Works with any version of R (2.11.1 or greater)
- Runs on Windows, Mac, Linux, and even over the web using RStudio Server
- Integrated with Git and Subversion for version control
- Free and open source (AGPLv3 license)

GUI ?

If you know EMACS then you will probably want to use ESS. See <http://ess.r-project.org/index.php?Section=home>



Emacs Speaks Statistics (ESS) provides an intelligent, consistent interface between the user and the software. ESS interfaces with SAS, S-PLUS, R, BUGS/JAGS and other statistical analysis packages on Unix, Linux and Microsoft Windows. ESS is itself a package within the emacs text editor and uses emacs features to streamline the creation and use of statistical software. ESS knows the syntax and grammar of statistical analysis packages and provides consistent display and editing features based on that knowledge. ESS assists in interactive and batch execution of statements written in these statistical analysis languages.

Finding Documentation for R

Stack Overflow: <http://stackoverflow.com>

Comprehensive R Archive Network: <http://cran.cnr.berkeley.edu/>

Manuals: <http://cran.cnr.berkeley.edu/manuals.html>

Good Intro Manual: http://cran.r-project.org/doc/contrib/Paradis-rdebuts_en.pdf

FAQs: <http://cran.cnr.berkeley.edu/faqs.html>

The R Journal: <http://journal.r-project.org/>

Contributed Documentation: <http://cran.cnr.berkeley.edu/other-docs.html>

BioConductor: <http://www.bioconductor.org> (400+ addons for life scientists)

R Graphical User Interface (GUI): http://sciviews.org/_rgui

R Programming Jobs: <http://www.programmingr.com/category/stype/r-job-listings>

Spreadsheet Addiction: http://www.burnsstat.com/pages/Tutor/spreadsheet_addiction.html

Finding Documentation for R

Learning Resources and Books

[edit]

There are many books on R as well as learning statistics with R. I have listed some free resources and some books that I have found to be useful. I would first start with the free resources since they are quite good and won't cost you anything to get started. If you are enrolled in a class then it is likely that your instructor will have some ideas, suggestions, and requirements for supporting R textbooks and learning resources.

Free resources (PDFs, websites, tutorials):

[edit]

- The R Inferno - Patrick Burns. Download the PDF http://www.burns-stat.com/pages/Tutor/R_inferno.pdf
- The R Programmer Wiki Book. See Wiki at http://en.wikibooks.org/wiki/R_Programming/
- Introduction to Probability and Statistics Using R - Jay Kerns. Download from <http://ipsur.org/index.html>
- Statistics with R - Vincent Zoonekynd. Website is at http://zoonek2.free.fr/UNIX/48_R/all.html
- Lattice Multivariate Data Visualization with R - Deepayan Sarkar <http://lmdvr.r-forge.r-project.org/figures/figures.html>
- Contributed R Information (A collection of free guides/handouts) <http://cran.r-project.org/other-docs.html>
- Rtips - Paul Johnson. <http://pj.freefaculty.org/R/Rtips.html#toc-Subsection-2.1>
- simpleR - Using R for Introductory Statistics - John Verzani - <http://www.unt.edu/rss/class/splus/Verzani-SimpleR.pdf>
- Do it yourself introduction to R - University of North Texas http://www.unt.edu/rss/class/Jon/R_SC/
- R Bloggers - A news aggregate site for R <http://www.r-bloggers.com/>
- R Journal - An open access, refereed journal of the R project for statistical computing. It features short to medium length articles covering topics that might be of interest to users or developers of R. <http://journal.r-project.org/index.html>

Finding Documentation for R

Books:

[\[edit\]](#)

Note that some of these books have freely available sample chapters on the Internet. If you check out these books on Amazon there is usually a "Look Inside" link so you can determine if the author's style appeals to your or not. Books can be quite expensive so be choosy when making your selection.

- R Cookbook - Paul Teetor
- R in a Nutshell - Joseph Adler
- The Art of R Programming: A Tour of Statistical Software Design - Norman Matloff
- Data Manipulation with R - Phil Spector
- ggplot2: Elegant Graphics for Data Analyses
- Introduction to Scientific Programming and Simulation Using R - Jones, Maillardet, Robinson
- BioConductor Case Studies - Hahne, Huber, Gentleman, Falcon
- Introductory Statistics with R - Peter Dalgaard
- Software for Data Analysis: Programming with R - John Chambers

Mailing Lists

[\[edit\]](#)

Here are some mailing lists that accept questions relative to R and BioConductor. Note that you should review the posting guidelines for each list before posting your question. You should also research the list to first see if your question has been answered. For newcomers to R this is almost always the case (that your question has already been answered previously). When posting - state your question/problem in specific terms along with a self-contained code example. Vague and open-ended questions are quite likely to remain unanswered.

- Cross-Validated <http://stats.stackexchange.com/> 
- R-Help <https://stat.ethz.ch/mailman/listinfo/r-help> 
- Stack Overflow <http://stackoverflow.com/questions/tagged/r> 

Finding Documentation for R

The Mailing Lists page contains general information and instructions for using the R-help mailing list. The general procedure is:

- 1) Subscribe to the R-help list at the Main R Mailing List (<https://stat.ethz.ch/mailman/listinfo/r-help>)
- 2) Read the Posting Guide (<http://www.r-project.org/posting-guide.html>) for information on writing an appropriate submission.
- 3) Before posting consider that many basic questions have been previously asked so first spend some time searching the list and its archives before posting.

About R-help

The main R mailing list, for announcements about the development of R and the availability of answers about problems and solutions using R, enhancements and patches to the source code, comparison and compatibility with S and S-plus, and for the posting of nice examples and best practices. See [General Instructions](#) on the [R Mailing Lists](#) page and follow the [posting guide](#)!

Finding Documentation for R

Style Guides for R Programming

Ever try reading code written by someone else ?



Take some "style" hints from the big guys: See "Google's R Style Guide" and Hadley's Style Guide at:

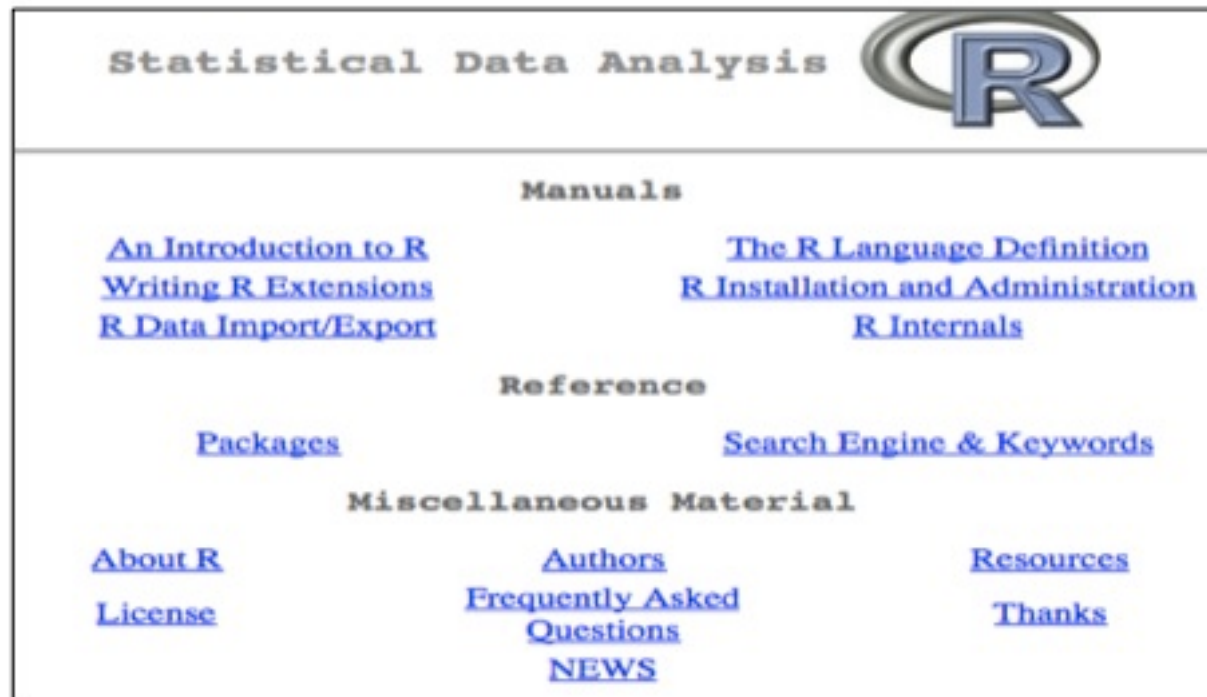
<http://google-styleguide.googlecode.com/svn/trunk/google-r-style.html>

<https://github.com/hadley/devtools/wiki/Style>

Workspace - Getting Help

R has built-in assistance that can be very useful. Let's take a look

```
> help.start()    # Launches a web browser with a table of contents.
```



Workspace - Getting Help

R has built-in assistance that can be very useful. Let's take a look

```
help(function_name)  # Get help on function
```

```
?function_name      # Equivalent to the above
```

```
args(function_name)  # See what arguments the function accepts
```

```
example(function_name)  # See an example
```

```
?mean
```

```
example(mean)
```

```
mean> x <- c(0:10, 50)
```

```
mean> xm <- mean(x)
```

```
mean> c(xm, mean(x, trim = 0.10))  
[1] 8.75 5.50
```

```
mean> mean(USArrests, trim = 0.2)  
Murder  Assault  UrbanPop  
7.42    167.60    66.20
```

Workspace - Getting Help

R has built-in assistance that can be very useful. Let's take a look

```
help.search("time series")
```

```
??"time series"
```

Help files with alias or concept or title matching 'time series' using fuzzy matching:

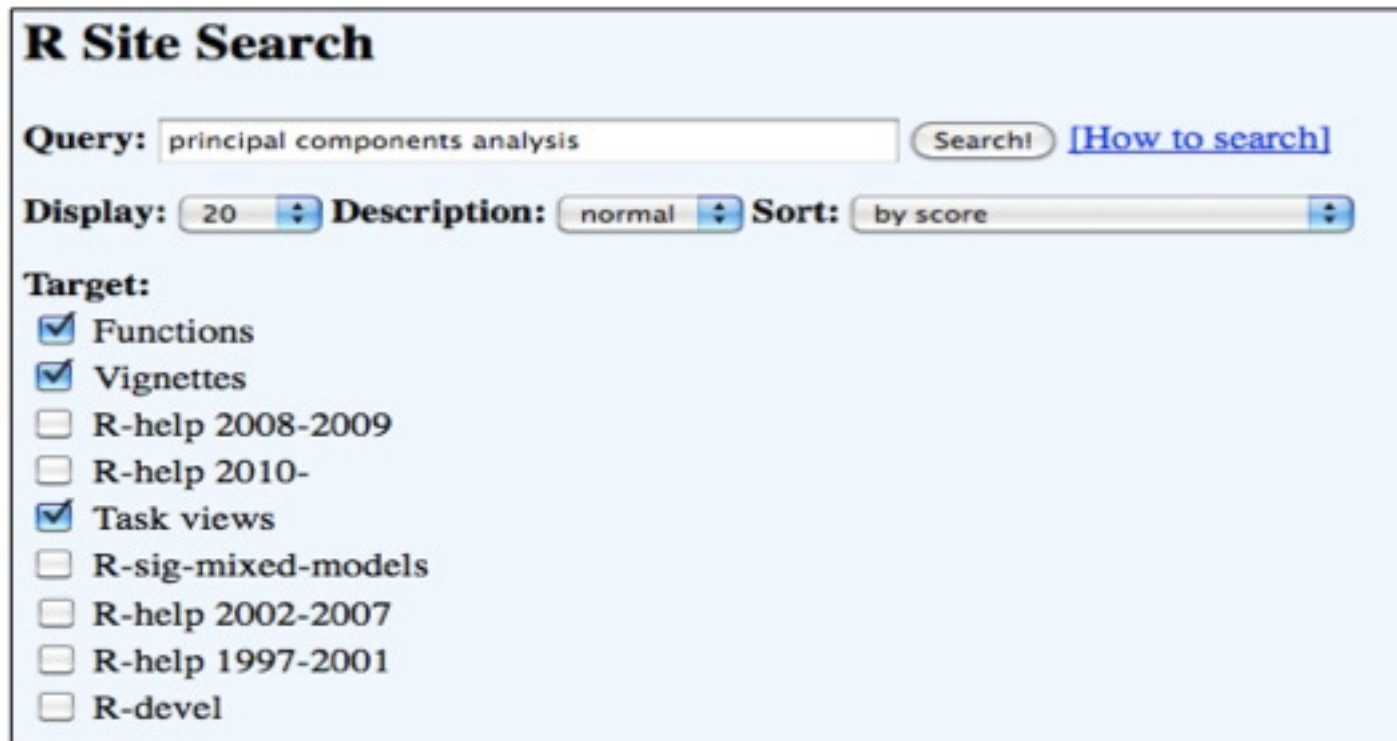
boot::tsboot	Bootstrapping of Time Series
car::Hartnagel	Canadian Crime-Rates Time Series
datasets::austres	Quarterly Time Series of the Number of Australian Residents
datasets::beavers	Body Temperature Series of Two Beavers
ggplot2::economics	US economic time series

Workspace - Getting Help

R has built-in assistance that can be very useful. Let's take a look

```
> RSiteSearch("correlation")
```

A search query has been submitted to <http://search.r-project.org>
The results page should open in your browser shortly



R Site Search

Query: [\[How to search\]](#)

Display: **Description:** **Sort:**

Target:

- ☒ Functions
- ☒ Vignettes
- ☐ R-help 2008-2009
- ☐ R-help 2010-
- ☒ Task views
- ☐ R-sig-mixed-models
- ☐ R-help 2002-2007
- ☐ R-help 1997-2001
- ☐ R-devel

R vs. SAS vs. MATLAB - Clash of the Titans



Let the debate begin..... But there are no easy answers here.
There are lots of good things about all three.

R

* Characteristics of R - R Journal, vol. 1/1 ([journal.r-project](http://journal.r-project.org))

1) Interactive language for data analysis and programming

2) Uses the functional programming model

3) Object-oriented

4) Modular - Accommodates add-on packages

5) Collaborative, Open source projects

* Everything that is in R is an object

* Everything that happens in R is a function call.

R

* The R main program is a classical read-eval-print loop

This basically means that you get a command prompt at which to enter things:

>

Three things to keep in mind:

- 1) The GREAT thing about R is that there are many different ways to do things.
- 2) The WORST thing about R is that there are many different ways to do things
- 3) To be a good programmer in R one must first be a knowledgeable user of R.

R - A Walkthrough

```
url = "http://www.bimcore.emory.edu/R/CSV.DIR/table_7_3.csv"
engine = read.table(url, sep = ",", header=TRUE)
```

```
engine=engine[,-1]
```

```
head(engine)          # 3 engine pollutants
```

```
   hc    co  nox
1 0.50  5.01 1.28
2 0.65 14.67 0.72
3 0.46  8.60 1.17
4 0.41  4.42 1.31
5 0.41  4.95 1.16
```

```
> summary(engine)
```

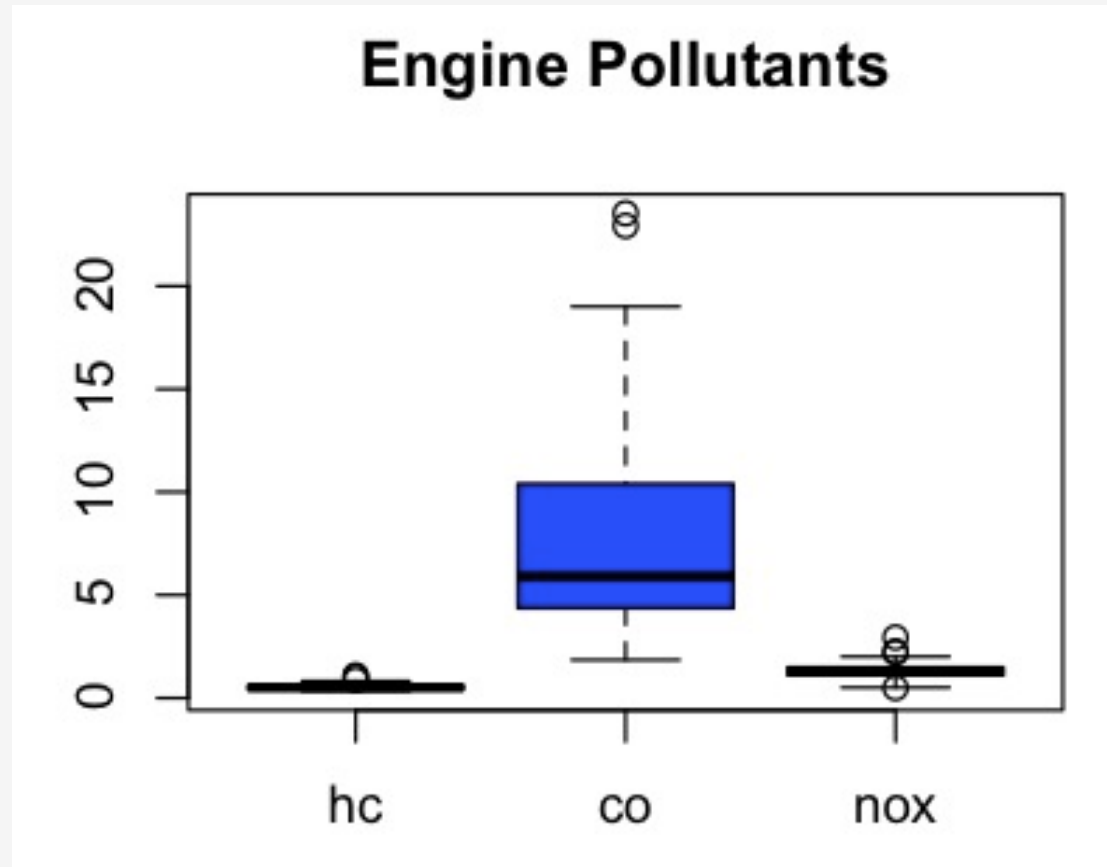
en	hc	co	nox
Min. : 1.00	Min. :0.3400	Min. : 1.850	Min. :0.490
1st Qu.:12.75	1st Qu.:0.4375	1st Qu.: 4.388	1st Qu.:1.110
Median :24.50	Median :0.5100	Median : 5.905	Median :1.315
Mean :24.00	Mean :0.5502	Mean : 7.879	Mean :1.340
3rd Qu.:35.25	3rd Qu.:0.6025	3rd Qu.:10.015	3rd Qu.:1.495
Max. :46.00	Max. :1.1000	Max. :23.530	Max. :2.940

<http://www.cyclismo.org/tutorial/R/cholesterol.html>

R - A Walkthrough

```
# Check out the data
```

```
boxplot(engine)
```



<http://www.cyclismo.org/tutorial/R/cholesterol.html>

R - A Walkthrough

```
# Plot some helpful graphs

par(mfrow=c(1,3))

boxplot(engine$co,main="Carbon Monoxide")

hist(engine$co)


qqnorm(engine$co,main="Carbon Monoxide")

qqline(engine$co)
```

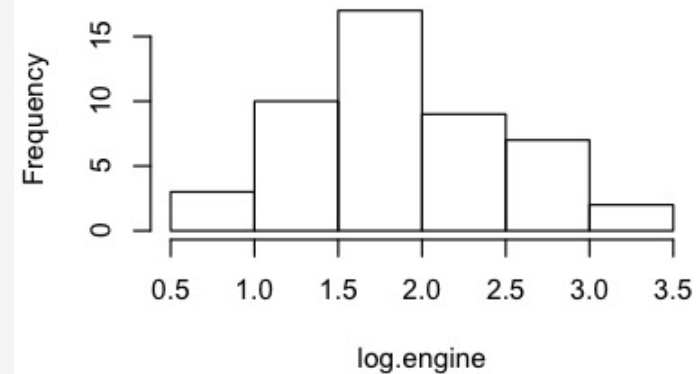
<http://www.cyclismo.org/tutorial/R/cholesterol.html>

R - A Walkthrough


Carbon Monoxide



Carbon Monoxide



QQ Plot for the Log of the Carbon Monoxide



R - A Walkthrough

```
# Null hypothesis is that the data is normal
```

```
shapiro.test(engine$co)
```

```
Shapiro-Wilk normality test
```

```
data: engine$co
```

```
W = 0.8357, p-value = 9.289e-06
```

```
# Take the log of the CO
```

```
log.engine = log(engine$co)
```

```
shapiro.test(log.engine)
```

```
Shapiro-Wilk normality test
```

```
data: log.engine
```

```
W = 0.9693, p-value = 0.2379
```

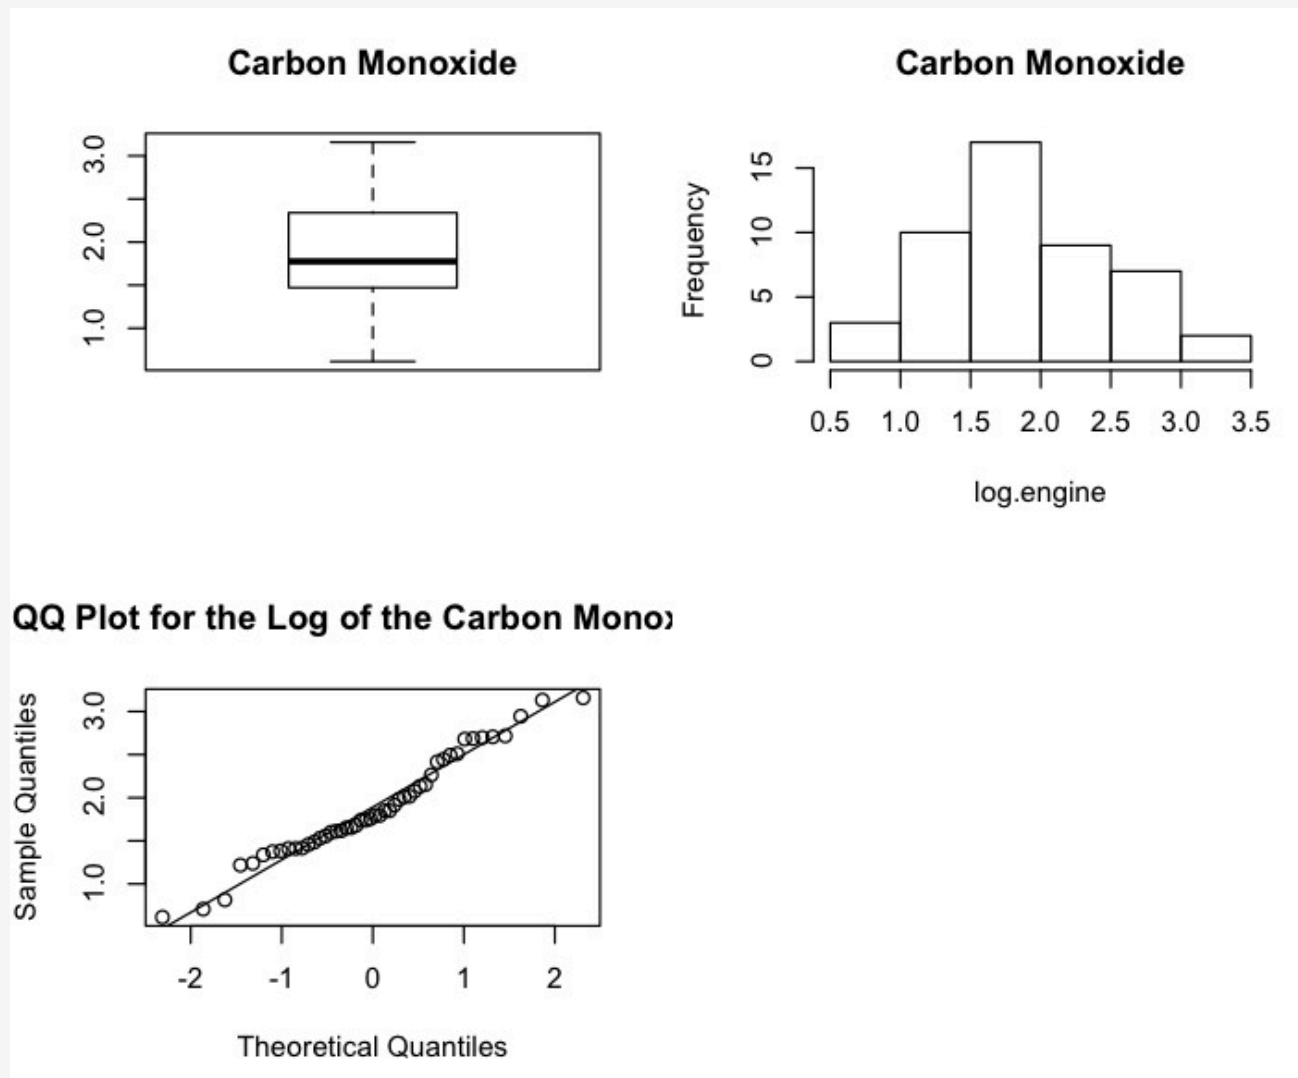
<http://www.cyclismo.org/tutorial/R/cholesterol.html>

R - A Walkthrough

```
par(mfrow=c(2,2))  
  
log.engine = log(engine$co)  
  
boxplot(log.engine,main="Carbon Monoxide")  
  
hist(log.engine,main="Carbon Monoxide")  
  
qqnorm(log.engine,main="QQ Plot for the Log of the Carbon Monoxide")  
  
qqline(log.engine)
```

<http://www.cyclismo.org/tutorial/R/cholesterol.html>

R - A Walkthrough



R - A Walkthrough

```
# Let's build a confidence interval
```

```
my.mean = mean(log.engine)
```

```
my.sd = sd(log.engine)
```

```
n = length(log.engine)
```

```
# Get standard error
```

```
se = my.sd/sqrt(n)
```

```
error = se*qt(0.975,df=n-1)
```

```
left = my.mean - error
```

```
right = my.mean + error
```

```
c(left,right)
```

```
[1] 1.709925 2.057431
```

```
c(exp(left),exp(right))
```

```
[1] 5.528548 7.825840
```

R - A Walkthrough

```
# Test H0: mu = 5.4
#      HA:mu != 5.4

lNull <- log(5.4) - error

rNull <- log(5.4) + error

c(lNull,rNull)
[1] 1.512646 1.860152

my.mean
[1] 1.883678
```

So the mean is outside the range so we reject the null hypothesis. There is a low probability that we would have obtained our sample mean if the true mean really were 5.4.

<http://www.cyclismo.org/tutorial/R/cholesterol.html>

R - A Walkthrough

```
# We could have calculated a p-value by hand
```

```
p.val = 2*(1-pt((my.mean-log(5.4))/se,df=n-1))
```

```
> p.val
```

```
[1] 0.02692539
```

```
# But its easier to call a procedure to do it all !!!!
```

```
t.test(log.engine,mu = log(5.4),alternative = "two.sided")
```

One Sample t-test

```
data: log.engine
```

```
t = 2.2841, df = 47, p-value = 0.02693
```

```
alternative hypothesis: true mean is not equal to 1.686399
```

```
95 percent confidence interval:
```

```
1.709925 2.057431
```

```
sample estimates:
```

```
mean of x
```

```
1.883678
```

<http://www.cyclismo.org/tutorial/R/cholesterol.html>

R - A Walkthrough

To find the power we need to set a level for the mean and then find the probability that we would accept the null hypothesis if the mean is really at the prescribed level. Here we will find the power to detect a difference if the level were 7.

```
power.t.test(n=n,delta=log(7)-log(5.4),sd=s,sig.level=0.05,  
             type="one.sample",alternative="two.sided",strict = TRUE)
```

One-sample t test power calculation

```
      n = 48  
    delta = 0.2595112  
      sd = 0.5983851  
sig.level = 0.05  
  power = 0.8371421  
alternative = two.sided
```

<http://www.cyclismo.org/tutorial/R/cholesterol.html>

R - Your First Real Session

```
?mean                # Get help on the mean function

example(kmeans)       # Run an example of kmeans (if it exists)

pi                    # Some popular quantities are built-in to R
[1] 3.141593

sqrt(2)              # Basic arithmetic
[1] 1.414214

print(pi)            # Print the comments of the pi variable
[1] 3.141593

X = 3; Y = 4         # Semicolon lets you enter 2 commands on the same line

Z = sqrt(X^2 + Y^2)   # Variables contain information

ls()                 # List all variables in the "environment"
[1] "X" "Y" "Z"
```

R - Your First Real Session - Calculator

```
2+3
[1] 5

3/2
[1] 1.5

2^3
[1] 8
4^2 - 3*2
[1] 10

(56-14)/6 - 4*7*10/(5^2-5)
[1] -7

abs(2-4)
[1] 2

cos(4*pi)
[1] 1

factorial(6)
[1] 720

choose(32,4)
[1] 35960

# Vector arithmetic is supported

x <- c(1,2,3,4)
y <- c(5,6,7,8)

x*y
[1] 5 12 21 32

y/x
[1] 5.000000 3.000000 2.333333 2.000000

y-x
[1] 4 4 4 4

cos(x*pi) + cos(y*pi)
[1] -2 2 -2 2
```

R - Your First Real Session - Calculator

```
log(10)
[1] 2.302585
```

```
log10(100)
[1] 2
```

```
sin(pi/2)
[1] 1
```

```
cos(pi/2)
[1] 6.123234e-17
```

```
> 1.3e6
[1] 1300000
```

```
> 9 %% 2
[1] 1
```

```
> 8 %% 2
[1] 0
```

```
> floor(5.7)
[1] 5
```

```
ceiling(6.8)
[1] 7
```

```
round(6.889,2)
[1] 6.89
```

```
3/0
[1] Inf
```

```
0/0
[1] NaN
```

```
is.finite(3)
[1]
```

```
x = c(1:8,NA)
x
[1] 1 2 3 4 5 6 7 8 NA
```

```
mean(x)
[1] NA
```

```
mean(x,na.rm=T)
[1] 4.5
```

R - Your First Real Session - Operators

Relational Operators

Equal to	==	if (myvar == "test") {print("EQ")}
	==	if (mynum == 3) {print("EQ")}
Not equal to	!=	if (myvar != "test") {print("NE")}
Less than or equal to	<=	if (number <= 5) {print("LTE")}
Less than	<	if (number < 10) {print("LT")}
Greater than or equal to	>=	if (number >= 10) {print("GTE")}
Greater than	>	if (number > 12) {print("GT")}

Boolean Operators

And	&	if ((myvar == "test") & (num <= 10)){ print("Equal and less than") }
Not	!	if (!complete.cases(myvec)) { print("Non complete cases") }
Or		if ((num > 3) (num < -3)) { print("Only one of these has to be true") }

R - Your First Real Session - Operators

Here are some popular math formulas rewritten in R. Note that many of these are implemented within R for convenient use. However it is helpful to gain a knowledge of operators and their precedence.

```
# a^2 + b^2 = c^2                                # Pythagorean Theorem
a = 2; b = 4
c = sqrt(a^2 + b^2)                                # To solve the PT for c
a = 2; b = 4; c = 1
(-b + sqrt(b^2 - 4*a*c)) / (2*a)                  # First case quadratic formula solution
(-b - sqrt(b^2 - 4*a*c)) / (2*a)                  # Second case quadratic formula solution
r = 4; h = 6; b = 3
circumference = 2*pi*r                            # circumference of a circle
area = (b*h)/2                                     # Area of a triangle
```


R - Your First Real Session - Operators

You can create expressions with R for later evaluation.

```
area = expression( (b*h)/2 )  
  
# Solve where b = 3 and h = 4  
  
b = 3  
  
h = 4  
  
eval(area)  
[1] 6
```

Note that you can easily make a function out of this also:

```
area <- function(b,h) {    # b and h are arguments / placeholders  
  my.area = (b*h)/2  
  return(my.area)  
}  
  
area(3,4)  
[1] 6
```

R - Your First Real Session - Operators

You can create expressions with R for later evaluation.

```
r1 = expression((-b + sqrt(b^2 - 4*a*c)) / (2*a))
```

```
r2 = expression((-b - sqrt(b^2 - 4*a*c)) / (2*a))
```

```
# Solve for  $ax^2 + bx + c$  where  $a = 1$ ,  $b=6$ , and  $c=8$ 
```

```
a = 1 ; b=6 ; c=8
```

```
eval(r1)  
[1] -2
```

```
eval(r2)  
[1] -4
```

```
a*eval(r1)^2 + b*eval(r1) + c  
[1] 0
```

```
a*eval(r2)^2 + b*eval(r2) + c  
[1] 0
```

R - Your First Real Session - Operators

We could also create a function out of this:

```
my.quad <- function(a,b,c) {  
  r1 = (-b + sqrt(b^2 - 4*a*c)) / (2*a)  
  r2 = (-b - sqrt(b^2 - 4*a*c)) / (2*a)  
  my.roots = c(r1,r2)  
  return(my.roots)  
}
```

Solve for $ax^2 + bx + c$ where $a = 1$, $b=6$, and $c=8$

```
my.quad(1,6,8)  
[1] -2 -4
```

Workspace - Startup

This is one of the more neglected areas when discussing R. However, it is important to understand how one gets information into and out of R. You can setup a file called `.Rprofile` in your home directory that enables you to configure default options for many things in R.

So if your home directory is `/home/wsp` or `/Users/fender` (assumes Linux or OSX) then create a file in your home directory called `.Rprofile`. Note that you are not obligated to use this capability but the more experience you get with R the more likely you will be interested in using it.

```
$ touch ~/.Rprofile  
$ vi .Rprofile
```

(Then edit it to contain things like on the next slide)

Workspace - Startup

```
# Things you might want to change

options(editor="notepad")
cd = setwd
pwd = getwd
lss = dir

# R interactive prompt
options(prompt="> ")
options(continue="+ ")

# General options
options(digits=3)
options(width = 130)
options(graphics.record=TRUE)

.First <- function(){
  library(Hmisc)
  cat("\nWelcome at", date(), "\n")
}

.Last <- function(){
  cat("\nGoodbye at ", date(), "\n")
}
```

Workspace - Navigating Directories and Files

This is one of the more neglected areas when discussing R. However, it is important to understand how one gets information into and out of R.

```
getwd()  
[1] "/Users/fender/TEST.DIR"
```

```
setwd("/Users/fender")
```

```
getwd()  
[1] "/Users/fender"
```

```
setwd("/Users/fender/TEST.DIR")
```

```
getwd()  
[1] "/Users/fender/TEST.DIR"
```

```
dir()  
[1] "coolpkg"          "coolpkg_1.0.tar.gz" "coolpkg.pdf"        "coolpkg.Rcheck"  
"g.Rd"              "stuff.R"
```

```
dir(,recursive=TRUE)  
[1] "coolpkg_1.0.tar.gz"          "coolpkg.pdf"  
[3] "coolpkg.Rcheck/00check.log"  "coolpkg.Rcheck/00install.out"  
[5] "coolpkg.Rcheck/coolpkg-Ex.pdf" "coolpkg.Rcheck/coolpkg-Ex.R"  
[7] "coolpkg.Rcheck/coolpkg-Ex.Rout" "coolpkg.Rcheck/coolpkg-manual.log"  
[9] "coolpkg.Rcheck/coolpkg-manual.pdf" "coolpkg.Rcheck/coolpkg/data/d.rda"
```

Workspace - Navigating Directories and Files

It is possible to manipulate files from within R. This doesn't mean that this is the best way to do this but it is a possibility.

```
R.home()
[1] "/Library/Frameworks/R.framework/Resources"

setwd(R.home())

list.files()
[1] "bin"          "COPYING"      "doc"          "etc"
     "fontconfig" "include"      "Info.plist"   "lib"
[9] "library"      "man1"         "modules"      "NEWS"
     "NEWS.pdf"    "R"            "Rscript"      "share"

list.files(path = "doc")
[1] "AUTHORS"      "COPYING"      "COPYRIGHTS"
     "CRAN_mirrors.csv" "FAQ"          "html"
[7] "KEYWORDS"     "KEYWORDS.db"  "manual"       "NEWS.rds"
"RESOURCES"
"THANKS"
```


Workspace -Navigating Directories and Files

It is possible to manipulate files from within R. This doesn't mean that this is the best way to do this but it is a possibility.

```
setwd("doc")
```

```
list.files()
```

```
[1] "AUTHORS"           "COPYING"           "COPYRIGHTS"
"CRAN_mirrors.csv" "FAQ"               "html"
[7] "KEYWORDS"          "KEYWORDS.db"       "manual"
"RESOURCES"        "THANKS"             "NEWS.rds"
```

```
path.expand("~")    # Finds you home directory in Linux
```

```
setwd(path.expand("~")) # Sets your working directory to be your home
```

Workspace - Navigating Directories and Files

It is possible to manipulate files from within R. This doesn't mean that this is the best way to do this but it is a possibility.

```
my.vec = list.files()
```

```
my.vec
[1] "AUTHORS"          "COPYING"          "COPYRIGHTS"       "CRAN_mirrors.csv"
"FAQ"               "html"
[7] "KEYWORDS"         "KEYWORDS.db"      "manual"           "NEWS.rds"
"RESOURCES"        "THANKS"
```

```
length(my.vec)
[1] 12
```

Workspace - Navigating Directories and Files

One can also use the “system” command to run commands at the UNIX level and reroute the output back into the R session.

```
system("ls -l")
```

```
AUTHORS
COPYING
COPYRIGHTS
CRAN_mirrors.csv
FAQ
KEYWORDS
KEYWORDS.db
NEWS.rds
RESOURCES
THANKS
html
manual
```

```
system("ls -l",intern=T)
```

```
[1] "AUTHORS"          "COPYING"          "COPYRIGHTS"
"CRAN_mirrors.csv" "FAQ"              "KEYWORDS"
[7] "KEYWORDS.db"      "NEWS.rds"         "RESOURCES"       "THANKS"
"html"              "manual"
```

Workspace - Saving Workspace

If you haven't noticed already, whenever you quit R it will prompt you to save your workspace. What this means is that you will have the ability to save a copy of all your current R objects (as evidenced by the `ls()` command). Generally you will want to save it.

Behind the scenes it stores the information into a BINARY file called `.Rdata` in whatever directory you are currently in.

```
> q()  
Save workspace image? [y/n/c]: y  
  
Goodbye at Mon Oct 1 14:26:47 2012  
  
fenders-macbook:TEST.DIR fender$ ls .Rdata  
.Rdata
```

When you start R back up in this directory it will load this file by default. A rookie mistake is when users start R from another directory expecting to find their objects reloaded

Workspace - Saving Workspace

If you want finer-grained control then you can specify your own file name.

```
save(list = ls(all=TRUE), file = "/Users/myhome/.MyRdata")
```

You can also elect to save only specific objects to a file.

```
my.lm = lm(mpg ~ wt,mtcars)
ls(my.lm)
[1] "assign"      "call"      "coefficients" "df.residual" "effects"
"fitted.values" "model"
[8] "qr"          "rank"      "residuals"   "terms"      "xlevels"

save(my.lm,file="/Users/myhome/.MyLM")
```

You can retrieve the contents of this file from any future R session by doing:

```
load("/Users/myhome/.MyLM")
```

Workspace - Saving Workspace

In general I suggest that you create a master directory to contain your R projects. You can then create sub-directories to house your various efforts.

That way when you come back to them after some lapse in time (e.g. semester break, vacation, etc) you can look at your directory structure and know where to pick back up.

```
$ ls RProjects
```

```
| -RProjects
  | -Genomes
    | ---1000_Genomes
    | ---Centenarians
  | -HIV
    | ---Replicates
  | -Influenza
```

Each subdirectory could have its own .RData file (or whatever you want to name it) to keep objects for that project intact. Don't pile up everything into one directory because it gets really confusing.

Workspace - Navigating

This is one of the more neglected areas when discussing R. However, it is important to understand how one gets information into and out of R.

```
ls()          # Check to see what objects there are in memory
character(0)
```

```
source("stuff.R",echo=TRUE)
```

```
f = function(x) x + y
g = function(x) x - y
```

```
ls()          # Ah, we now have two functions in the workspace.
[1] "f" "g"
```


Workspace - Navigating

This is one of the more neglected areas when discussing R. However, it is important to understand how one gets information into and out of R.

```
ls()          # Ah, we no have two functions in the workspace.  
[1] "f" "g"
```

```
rm("g")
```

```
ls()  
[1] "f"
```

```
rm(list = ls())  # Removes all objects in the workspace  
                 # This can be dangerous - use with care
```

Workspace - Basic I/O - Variables

As in most programming languages, it is customary to store or hold the results of an operation in a variable name. In R such results are assigned with the symbols "<-" or "=". Variable names are case sensitive.

```
X <- 2.5    # These two statements are equivalent
```

```
X = 2.5 # Same as above
```

```
X  
[1] 2.5
```

```
mynewvar = X + 3
```

```
MYNEWVAR = X + 3    # Two different variables
```

R has several one-letter reserved words: **c**, **q**, **s**, **t**, **C**, **D**, **F**, **I**, and **T**

You cannot begin a variable name with a period character "."

Workspace - Basic I/O - Variables

Valid Names

Variable names in R are case-sensitive, so myvar is not the same as MYVAR

Variables names should not begin with numbers (e.g. 1) or symbols (%,_)

Variable names should not contain spaces in the name (my var)

Workspace - Basic I/O - Variables

Valid Variable Names

```
mean.height  
smoker  
non.smoker  
temp.var  
patient_id  
Eye.Color  
State_Population  
disease.state  
White_Cell_Count  
jobTitle
```

Invalid Variable Names

```
.mean.height  
_myvariable  
_Mean.height  
lvariable  
l_variable  
%some.var  
some var  
"some var"
```

Scope

R uses lexical scoping, which we will explore in some detail later.

Since variables cannot be declared, (they pop into existence on first assignment) it is not always easy to determine the scope of a variable. You cannot tell just by looking at the source code of a function whether a variable is local to that function.

Workspace - Basic I/O

You can use the scan function to enter data from the command line. This is also useful when pasting in a row of numbers from a website.

```
my.input.vector = scan()  
1: 8 9 10 11 23 48 73  
8:  
Read 7 items
```

```
my.input.vector  
[1] 8 9 10 11 23 48 73
```

The cat command can be used to take some input and write it to a file:

```
cat("TITLE extra line", "2 3 5 7", file="ex.data")  
cat("TITLE extra line", "2 3 5 7", "11 13 17",  
+ file="ex.data", sep="\n")  
  
dir(pattern="ex.data")  
[1] "ex.data"
```

Workspace - Basic I/O

You can prompt users for input but this isn't really all that common. But let's say you wanted to prompt the user to enter a number:

```
n <- readline("enter a positive integer: ")
enter a positive integer: 23
```

```
n
[1] "23"
```

You could make this more involved by doing some error checking to make sure they enter in an integer.

```
n = -1

while(n < 1 ){
  n <- readline("enter a positive integer: ")
  n <- ifelse(grepl("\\D",n),-1,as.integer(n))
  if(is.na(n)){break} # breaks when hit enter
}
enter a positive integer: sdf5
enter a positive integer: 45
```

Workspace - Basic I/O

It is also very easy to turn around and write data frames and matrices to .CSV file for use with any program that uses .CSV files (e.g. SAS, SPSS, Minitab, Excel, etc).

```
head(mtcars)
```

	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
Mazda RX4	21.0	6	160	110	3.90	2.620	16.46	0	1	4	4
Mazda RX4 Wag	21.0	6	160	110	3.90	2.875	17.02	0	1	4	4
Datsun 710	22.8	4	108	93	3.85	2.320	18.61	1	1	4	1
Hornet 4 Drive	21.4	6	258	110	3.08	3.215	19.44	1	0	3	1

```
write.table(mtcars,file="mtcars.csv",  
            row.names=TRUE,           # Row names get saved  
            col.names=TRUE,          # Header gets saved  
            sep=",")
```

```
$ head mtcars.csv  
"mpg","cyl","disp","hp","drat","wt","qsec","vs","am","gear","carb"  
"Mazda RX4",21,6,160,110,3.9,2.62,16.46,0,1,4,4  
"Mazda RX4 Wag",21,6,160,110,3.9,2.875,17.02,0,1,4,4
```


Workspace - CSV files

You can also read text files a line at a time use the `readLines()` function. Again, this is not so common an activity but it is something you should have some exposure to. Let's say we have a comma delimited file.

Workspace - "Sinking" your work

From time to time you might want to capture the output of your work though as you already know this is possible using the "save" command to save the contents of your environment. The "cat", "write", and "write.table" commands allow you to save specific variables to a file.

The "sink" command exists to provide a way to redirect the output from R commands directly to a file. This might be something you would do if you were running a very long running batch job. This would let you see the progress of your program on an incremental basis.

Let's say we have the following R statements:

```
set.seed(123)
x <- rnorm(10)
y <- rnorm(10)

print(x)
cat ("y =", y, "\n")

t.test(x,y)
plot(x,y)
```

Workspace - "Sinking" your work

If we run this interactively the output from this would look like the following. We would also see a plot window because of the call to the plot command.

```
set.seed(123)
x <- rnorm(10)
y <- rnorm(10)

print(x)
[1] -0.56047565 -0.23017749  1.55870831  0.07050839  0.12928774  1.71506499
[7]  0.46091621 -1.26506123 -0.68685285 -0.44566197

cat ("y =", y, "\n")
y = 1.224082 0.3598138 0.4007715 0.1106827 -0.5558411 1.786913 0.4978505 -1.966617
0.7013559 -0.4727914

t.test(x,y)
    Welch Two Sample t-test

data:  x and y
t = -0.3006, df = 17.872, p-value = 0.7672
alternative hypothesis: true difference in means is not equal to 0
95 percent confidence interval:
 -1.0710488  0.8030562
sample estimates:
 mean of x  mean of y
0.07462564 0.20862196
```

Workspace - "Sinking" your work

If we wanted we could redirect all the output from the print,cat statement and t.test function to a file called "my.results.txt"

```
sink("my.results.txt") # All output will now go to "my.results.txt"
```

```
set.seed(123)
```

```
x <- rnorm(10)
```

```
y <- rnorm(10)
```

```
print(x)
```

```
cat ("y =", y, "\n")
```

```
t.test(x,y)
```

```
plot(x,y)
```

```
sink() # This will turn off the sink
```

If we run this then we see only the statements as they are processed but not any of the output. To see the output we look at "my.results.txt"

Workspace - "Sinking" your work

If we run this then we see only the statements as they are processed but not any of the output. To see the output we look at "my.results.txt". Note that the result of the plot command will go by default to a file named "Rplots.pdf" located in the current working directory.

```
$ more my.results.txt
```

```
[1] -0.56047565 -0.23017749  1.55870831  0.07050839  0.12928774  1.71506499
[7]  0.46091621 -1.26506123 -0.68685285 -0.44566197
y = 1.224082 0.3598138 0.4007715 0.1106827 -0.5558411 1.786913 0.4978505
-1.966617 0.7013559 -0.4727914
```

```
Welch Two Sample t-test
```

```
data:  x and y
t = -0.3006, df = 17.872, p-value = 0.7672
alternative hypothesis: true difference in means is not equal to 0
95 percent confidence interval:
 -1.0710488  0.8030562
sample estimates:
 mean of x  mean of y
0.07462564 0.20862196
```

Workspace - "Sinking" your work

Okay, an alternative approach to using sink is to run your R jobs as "batch" jobs. Let's put these commands into a file called test.R. Note the we don't have any "sink" commands.

```
set.seed(123)
x <- rnorm(10)
y <- rnorm(10)

print(x)
cat ("y =", y, "\n")

t.test(x,y)
plot(x,y)
```

We can run this as a batch job on Linux and OSX like:

```
$ R CMD BATCH test.R myout
```

This will put your results into a file "myout". As before the plots will be stashed into a file called "Rplots.pdf".

Workspace - "Sinking" your work

If you wanted more control over how your plot was named then you can use one of the functions designed to create plots in a known format (PNG, JPEG, PDF).

```
set.seed(123)
x <- rnorm(10)
y <- rnorm(10)

print(x)
cat ("y =", y, "\n")

t.test(x,y)

pdf("myplots.pdf")    # Redirects plots to myplots.pdf

plot(x,y)

dev.off()             # Turns off plot redirection
```

We can run this as a batch job on Linux and OSX like:

```
$ R CMD BATCH test.R myout
```

Workspace - Tables

Note that you can read CSV files directly from the Internet as long as you have the URL. AND there are dedicated functions to intelligently parse the data so you don't have to do all the things you had to do in the previous slide.

```
url = "http://www.bimcore.emory.edu/BIOS560R/DATA.DIR/hsb2.csv"
```

```
my.input = read.table(url,header=T,sep=",")
```

```
head(my.input)
```

	gender	id	race	ses	schtyp	prgtype	read	write	math	science	socst
1	0	70	4	1	1	general	57	52	41	47	57
2	1	121	4	2	1	vocati	68	59	53	63	61
3	0	86	4	3	1	general	44	33	54	58	31
4	0	141	4	3	1	vocati	63	44	47	53	56
5	0	172	4	2	1	academic	47	52	57	53	61
6	0	113	4	2	1	academic	44	52	51	63	61

Workspace - Tables

Sometimes you want to read data from the Internet but its not in a CSV file. Its in a HTML table. Take a look at this example. Let's say we want to get the data in the first table from here: <http://msenux.redwoods.edu/math/R/regression.php>

Scatterplots

The data set in the table that follows is taken from measured the heights of 161 children in Kalama, a village recorded each month, with the study lasting several years follows.

Mean Height versus Age	
Age in Months	Average Height in Centimeters
18	76.1
19	77
20	78.1
21	78.2
22	78.8
23	79.7

Workspace - Tables

Let's write a little R code to get this information:

```
library(XML)
url = "http://msenux.redwoods.edu/math/R/regression.php"
my.table = readHTMLTable(url,which=1)
```

```
my.table
```

	V1	V2
1	Mean Height versus Age	<NA>
2	Age in Months	Average Height in Centimeters
3	18	76.1
4	19	77
..		
..		

```
my.table = my.table[-1:-2,]
names(my.table) = c("Age","Height")
head(my.table)
```

	Age	Height
3	18	76.1
4	19	77
5	20	78.1
6	21	78.2
7	22	78.8

Workspace - Complicated Tables

Let's look at a more involved example. Look at the follow info on World Population http://en.wikipedia.org/wiki/World_population

This page has many tables. Let's get the data corresponding to the one that looks like:

The 10 countries with the largest total population:

Rank	Country / Territory	Population	Date	% of world population	Source
1	 China ^[note 2]	1,353,430,000	October 1, 2012	19.2%	[71]
2	 India	1,210,193,422	March 2011	17%	[72]
3	 United States	314,490,000	October 1, 2012	4.47%	[73]
4	 Indonesia	238,400,000	May 2010	3.33%	[74]
5	 Brazil	197,067,000	October 1, 2012	2.8%	[75]
6	 Pakistan	180,819,000	October 1, 2012	2.57%	[76]
7	 Nigeria	170,123,740	July 2012	2.42%	[77]
8	 Bangladesh	161,083,804	July 2012	2.29%	[78]
9	 Russia	141,927,297	January 1, 2010	2.015%	[79]
10	 Japan	127,610,000	May 1, 2012	1.81%	[80]

Workspace - Complicated Tables

This appears to be the 5th table on the page. Note that this is one of those things that you will probably have to try a few times before you get it right.

```
url = "http://en.wikipedia.org/wiki/World\_population"
```

```
table.four = readHTMLTable(url, which=4) # Most densely populated countries
```

```
table.four
```

	Rank	Country / Territory	Population	Date	% of world\npopulation	Source
1	1	China[note 2]	1,353,430,000	October 1, 2012	19.2%	[71]
2	2	India	1,210,193,422	March 2011	17%	[72]
3	3	United States	314,490,000	October 1, 2012	4.47%	[73]
4	4	Indonesia	238,400,000	May 2010	3.33%	[74]

Okay, we got the data but its a little messy. Let's clean it up some.

Workspace - Complicated Tables

Okay, we got the data but its a little messy. Let's clean it up some.

```
table.four = table.four[,-4:-6] # Eliminate the 4th-6th column
table.four
```

	Rank	Country / Territory	Population
1	1	China[73]	1,352,190,000
2	2	India	1,210,193,422

```
# Get rid of the commas in the numbers
```

```
table.four$Population=as.numeric(gsub(",","",table.four$Population))/100000
```


```
# Give the columns new names
```

```
names(table.four) = c("Rank","Country","Population")
```

```
# Plot the data
library(lattice)
```

```
xyplot(Population ~ Country,table.four,scales = list(x = c(rot=60)),
       type="h",main="Most Densely Populated Countries")
```

Workspace - Complicated Tables



Workspace - Other Stats Packages

Its also possible to read data sets from other statistical packages. However, as with Excel, I suggest you export those datasets into a CSV format and then import them into R.

Function(s)	Purpose
<code>read.epinfo</code>	Read saved objects from EpiInfo
<code>read.xport</code>	Read saved objects in SAS export format
<code>read.spss</code>	Read saved objects from SPSS written using the save or export command
<code>read.systat</code>	Read saved objects from SYSTAT rectangular (mtype=1) data only
<code>read.dta</code>	Read saved objects from STATA (versions 5-9)
<code>read.mtp</code>	Read Minitab Portable Worksheet Files
<code>read.octave</code>	Read saved objects from GNU octave
<code>read.dbf</code>	Read or write saved objects from DBF files (FoxPro, dBase,etc)

Workspace - Excel

You can even read directly from Excel spreadsheets but this is usually something you wouldn't do. Always try to save the data out of Excel into a CSV file and then import it into R.

If you simply MUST read the Excel spreadsheet directly then use an approach similar to the following. Note that I've never done this so can't say that it will work. Also, this might work only on Windows machines.

```
> library(RODBC)
> channel <- odbcConnectExcel("examp.xls")
## list the spreadsheets
> sqlTables(channel)
  TABLE_CAT TABLE_SCHEM      TABLE_NAME  TABLE_TYPE REMARKS
1 C:\\bdr      NA          Sheet1$  SYSTEM TABLE      NA
2 C:\\bdr      NA          Sheet2$  SYSTEM TABLE      NA
3 C:\\bdr      NA          Sheet3$  SYSTEM TABLE      NA
4 C:\\bdr      NA Sheet1$Print_Area      TABLE      NA

## retrieve the contents of sheet 1, by either of
> sh1 <- sqlFetch(channel, "Sheet1")
> sh1 <- sqlQuery(channel, "select * from [Sheet1$]")
```


Workspace - Databases

Note that you can also read and write to relational databases. We will explore this much more in depth in a later module. You may wonder - Why should you care about connecting to relational databases from R. Good reasons include:

- * There are limitations on the types of data that R handles well. Since all data being manipulated by R are resident in memory, and several copies of the data can be created during execution of a function,
- * R is not well suited to extremely large data sets. Data objects that are more than a (few) hundred megabytes in size can cause R to run out of memory, particularly on a 32-bit operating system.
- * R does not easily support concurrent access to data. That is, if more than one user is accessing, and perhaps updating, the same data, the changes made by one user will not be visible to the others.
- * R does support persistence of data, in that you can save a data object or an entire worksheet from one session and restore it at the subsequent session, but the format of the stored data is specific to R and not easily manipulated by other systems.

Workspace - Databases

Database management systems (DBMSs) and, in particular, relational DBMSs (RDBMSs) *are* designed to do all of these things well. Their strengths are:

1. To provide fast access to selected parts of large databases.
2. Powerful ways to summarize and cross-tabulate columns in databases.
3. Store data in more organized ways than the rectangular grid model of spreadsheets and R data frames.
4. Concurrent access from multiple clients running on multiple hosts while enforcing security constraints on access to the data.
5. Ability to act as a server to a wide range of clients.

The sort of statistical applications for which DBMS might be used are to extract a 10% sample of the data, to cross-tabulate data to produce a multi-dimensional contingency table, and to extract data group by group from a database for separate analysis.

Workspace - Databases

Note that you can also read and write to relational databases. We will explore this much more in depth in a later module. For now here is an example of how you might connect to a MySQL database.

```
m <- dbDriver("SQLite")

# initialize a new database to a tempfile and copy some data.frame
# from the base package into it

tfile <- tempfile()
con <- dbConnect(m, dbname = tfile)
data(USArrests)
dbWriteTable(con, "USArrests", USArrests)

# query
rs <- dbSendQuery(con, "select * from USArrests")
d1 <- fetch(rs, n = 10)      # extract data in chunks of 10 rows
d1
```

	row_names	Murder	Assault	UrbanPop	Rape
1	Alabama	13.2	236	58	21.2
2	Alaska	10.0	263	48	44.5
3	Arizona	8.1	294	80	31.0
4	Arkansas	8.8	190	50	19.5
5	California	9.0	276	91	40.6
6	Colorado	7.9	204	78	38.7

Workspace - Databases

Note that you can also read and write to relational databases. We will explore this much more in depth in a later module. For now here is an example of how you might connect to a MySQL database.

```
## Select from the loaded table
```

```
dbClearResult(rs)
```

```
rs <- dbSendQuery(con, "select * from USArrests where  
                        Assault > 10 order by Murder")
```

```
d1 <- fetch(rs, n = 10)
```

```
d1
```

	row_names	Murder	Assault	UrbanPop	Rape
1	North Dakota	0.8	45	44	7.3
2	Maine	2.1	83	51	7.8
3	New Hampshire	2.1	57	56	9.5
4	Iowa	2.2	56	57	11.3
5	Vermont	2.2	48	32	11.2
6	Idaho	2.6	120	54	14.2
7	Wisconsin	2.6	53	66	10.8
8	Minnesota	2.7	72	66	14.9
9	Utah	3.2	120	80	22.9
10	Connecticut	3.3	110	77	11.1

Workspace - Databases

Another way to work with data sets as if they were databases is to use the "sqldf" package. We'll look more closely at this in the data frame section. This package is very useful if you already know SQL and are just learning R. sqldf is an add on package so you will first need to install it.

```
install.packages("sqldf", dependencies = TRUE)

library(sqldf)

data(mtcars)

sqldf("select * from mtcars where mpg > 20 AND cyl == 4")
```

Workspace - XML

R can also process XML files from the Internet, which is very powerful. As an example we'll access Google's GeoCoding pages to get the Latitude and Longitude for Atlanta, GA from within a program. Its just a simple example.

<https://developers.google.com/maps/documentation/geocoding/>



Workspace - XML

R can also process XML files from the Internet, which is very powerful. As an example we'll access Google's GeoCoding pages to get the Latitude and Longitude for Atlanta, GA from within a program. Its just a simple example.

```
library(RCurl)
library(XML)

my.url = "http://maps.googleapis.com/maps/api/geocode/xml?
address=Atlanta,GA&sensor=false"

txt = getURL(my.url)
hold = xmlTreeParse(txt,useInternalNodes=TRUE)

hold
<?xml version="1.0" encoding="UTF-8"?>
<GeocodeResponse>
  <status>OK</status>
  <result>
    <type>locality</type>

place = getNodeSet(hold,"//GeocodeResponse/result[1]/geometry/location[1]/*")
as.numeric(sapply(place,xmlValue))
[1] 33.74900 -84.38798
```

Workspace - XML

R can also process XML files from the Internet, which is very powerful. As an example we'll access Google's GeoCoding pages to get the Latitude and Longitude for Atlanta, GA from within a program. Its just a simple example.

```
- <GeocodeResponse>
  <status>OK</status>
  - <result>
    <type>locality</type>
    <type>political</type>
    <formatted_address>Atlanta, GA, USA</formatted_address>
  - <address_component>
    <long_name>Atlanta</long_name>
    <short_name>Atlanta</short_name>
    <type>locality</type>
    <type>political</type>
  </address_component>
  - <address_component>
    <long_name>Fulton</long_name>
    <short_name>Fulton</short_name>
    <type>administrative_area_level_2</type>
    <type>political</type>
  </address_component>
```


Workspace - XML

The interesting thing is that we can then use this information to query other resources. Say for example we can look at the Personal Weather Stations Web site and get quick assessment of the weather in the area based on a survey of weather stations.

```
my.url = "http://api.wunderground.com/auto/wui/geo/GeoLookupXML/index.xml?
query=33.749,-84.38798"
```

```
myweather = getURL(url)
hold = xmlTreeParse(myweather, useInternalNodes=TRUE)
```

```
qstr = "//location/nearby_weather_stations/pws/station[distance_mi<5]/
ancestor-or-self::station/id"
```

```
stationsXml = getNodeSet(hold,qstr)
```

```
stations = sapply(stationsXml,xmlValue)
```

```
> stations
[1] "KGAATLAN40" "KGAATLAN49" "KGAATLAN37" "KGAATLAN68" "KGAATLAN57"
[6] "KGAATLAN54" "KGAATLAN16"
```